

HW 7 : GANs

In this assignment, you will learn to write an advanced PyTorch implementation concept commonly used in a complex deep learning pipeline by using GAN as a learning example.

You will also start working with more complex architectures (upsampling) and writing style (complex modules such as modulelist) for pytorch.

Every TODO is weighted equally. Optional TODO is half of a regular TODO.

GPU test

```
!nvidia-smi

Thu Apr  4 19:44:00 2024
+-----+
| NVIDIA-SMI 535.129.03      Driver Version: 535.129.03 CUDA
Version: 12.2                |
+-----+
| GPU  Name                  Persistence-M | Bus-Id        Disp.A |
| Volatile Uncorr. ECC |          |
| Fan  Temp    Perf          Pwr:Usage/Cap |         Memory-Usage |
| GPU-Util  Compute M. |          |
| MIG M.   |          |
|          |          |
+-----+
| 0  Tesla T4               Off  | 00000000:00:04.0 Off |
| 0 |                                |
| N/A   35C     P8              9W / 70W |       0MiB / 15360MiB |
| 0%   Default |          |
|          |          |
| N/A |          |
+-----+
| 1  Tesla T4               Off  | 00000000:00:05.0 Off |
| 0 |                                |
| N/A   35C     P8              9W / 70W |       0MiB / 15360MiB |
| 0%   Default |          |
|          |          |
| N/A |          |
+-----+
```

Processes:					
GPU	GI	CI	PID	Type	Process name
GPU Memory					
	ID	ID			
Usage					
=====					
No running processes found					
=====					

Part 1 : WGAN-GP reimplementaion

In this section, you are going to reimplement WGAN-GP (<https://arxiv.org/pdf/1704.00028.pdf>) based on the pseudocode provided in the paper to generate MNIST digit characters. Some parts are intentionally modified to discourage straight copy/pasting from public repositories.

Algorithm 1 WGAN with gradient penalty. We use default values of $\lambda = 10$, $n_{\text{critic}} = 5$, $\alpha = 0.0001$, $\beta_1 = 0$, $\beta_2 = 0.9$.

Require: The gradient penalty coefficient λ , the number of critic iterations per generator iteration n_{critic} , the batch size m , Adam hyperparameters α, β_1, β_2 .

Require: initial critic parameters w_0 , initial generator parameters θ_0 .

```

1: while  $\theta$  has not converged do
2:   for  $t = 1, \dots, n_{\text{critic}}$  do
3:     for  $i = 1, \dots, m$  do
4:       Sample real data  $\mathbf{x} \sim \mathbb{P}_r$ , latent variable  $\mathbf{z} \sim p(\mathbf{z})$ , a random number  $\epsilon \sim U[0, 1]$ .
5:        $\hat{\mathbf{x}} \leftarrow G_\theta(\mathbf{z})$ 
6:        $\hat{\mathbf{x}} \leftarrow \epsilon \mathbf{x} + (1 - \epsilon) \hat{\mathbf{x}}$ 
7:        $L^{(i)} \leftarrow D_w(\hat{\mathbf{x}}) - D_w(\mathbf{x}) + \lambda (\|\nabla_{\hat{\mathbf{x}}} D_w(\hat{\mathbf{x}})\|_2 - 1)^2$ 
8:     end for
9:      $w \leftarrow \text{Adam}(\nabla_w \frac{1}{m} \sum_{i=1}^m L^{(i)}, w, \alpha, \beta_1, \beta_2)$ 
10:   end for
11:   Sample a batch of latent variables  $\{\mathbf{z}^{(i)}\}_{i=1}^m \sim p(\mathbf{z})$ .
12:    $\theta \leftarrow \text{Adam}(\nabla_\theta \frac{1}{m} \sum_{i=1}^m -D_w(G_\theta(\mathbf{z})), \theta, \alpha, \beta_1, \beta_2)$ 
13: end while

```

The pseudocode could be organized into two main parts: discriminator optimization in line 2-10, and generator optimization in line 11-12.

The discriminator part consists of four steps:

- Line 4: data, and noise sampling with a batch size of m
- Line 5-7: discriminator loss calculation

- Line 9: discriminator update
- Repeat line 4-9 for n_{critic} steps

After the discriminator is updated, the generator is then updated by performing two steps:

- Line 11: noise sampling
- Line 12: generator loss calculation and update

This part is divided into four subsections: network initialization, hyperparameter initialization, data preparation, and training loop. The detail for each part will be explained in the subsections.

Downloading MNIST dataset

The MNIST dataset contains 60,000 training digit character image (0-9) at 28x28 resolution that are normalized to [0, 1]. Given the training images, your task is to generate new training images using WGAN-GP by learning from the training distribution.

```
import torchvision.datasets as datasets
import numpy as np

mnist_trainset = datasets.MNIST(root='./data', train=True,
download=True, transform=None)
trainX = np.array(mnist_trainset.data[..., None]).transpose(0, 3, 1,
2) / 255
print("Dataset size : ", trainX.shape)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz

100%|██████████| 9912422/9912422 [00:00<00:00, 103999597.88it/s]

Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to
./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz

100%|██████████| 28881/28881 [00:00<00:00, 43953444.78it/s]

Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to
./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
```

```

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz

100%|██████████| 1648877/1648877 [00:00<00:00, 27640568.63it/s]

Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to
./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz

100%|██████████| 4542/4542 [00:00<00:00, 8384915.83it/s]

Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to
./data/MNIST/raw

Dataset size : (60000, 1, 28, 28)

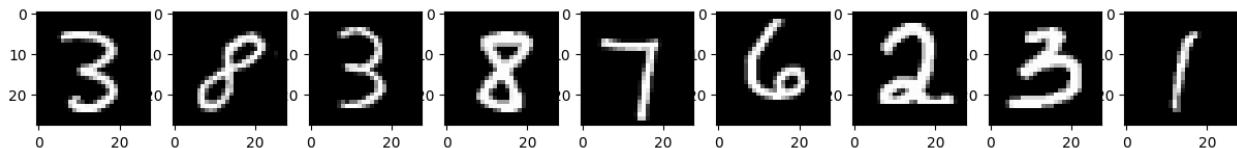
```

Dataset Visualization

```

import matplotlib.pyplot as plt
import numpy as np
plt.figure(figsize = (15,75))
for i in range(9):
    plt.subplot( int('19{}'.format(i+1)) )
    plt.imshow( trainX[np.random.randint(len(trainX)).transpose((1, 2, 0))[:, :, 0] , cmap = 'gray' ) )
plt.show()

```



Generator and Discriminator network

Before training, the deep learning networks have to be initialized first. Therefore, in this part, you are going to write a generator and discriminator network based on the description provided below.

The description of the discriminator network is shown in the Table below.

Discriminator (D(x))			
	Kernel size	Resample	Output shape
ConvBlock	5×5	Down	$128 \times 14 \times 14$
ConvBlock	5×5	Down	$256 \times 7 \times 7$
ConvBlock	5×5	Down	$512 \times 4 \times 4$
Linear	-	-	1

The network also has some specific requirements:

- ConvBlock is a Convolution-ReLU layer
- All ReLUs in the encoder are leaky, with a slope of 0.1

The description of the generator network is shown in the Table below.

Generator (G(z))			
	Kernel size	Resample	Output shape
z	-	-	128
Linear	-	-	$512 \times 4 \times 4$
ConvBlock	5×5	Up	$256 \times 8 \times 8$
ConvBlock	5×5	Up	$128 \times 16 \times 16$
Conv, Sigmoid	5×5	Up	$1 \times 32 \times 32$
-	-	Down	$1 \times 28 \times 28$

The network also has some specific requirements:

- ConvBlock is a ConvTranspose-BatchNorm-ReLU layer
- Downsampling method is bilinear interpolation (torch.nn.Upsample or torch.nn.functional.interpolate)

TODO 1: Implement a discriminator network. TODO 2: Implement a generator network.

```
!pip install torch-summary
Collecting torch-summary
  Downloading torch_summary-1.4.5-py3-none-any.whl.metadata (18 kB)
  Downloading torch_summary-1.4.5-py3-none-any.whl (16 kB)
Installing collected packages: torch-summary
Successfully installed torch-summary-1.4.5
```

```

import torch
import torch.nn.functional as F
from torch import nn
from torchvision import transforms
from torchsummary import summary

def convBlock(conv_config=None, norm_config=None, relu_config=None,
dropout_config=None, sigmoid_config=None):
    block = nn.Sequential()
    if conv_config != None:
        cv_name, cv_type, in_channels, out_channels, kernel_size,
stride, padding, output_padding = conv_config
        if cv_type == "default":
            block.add_module(cv_name, nn.Conv2d(in_channels,
out_channels, kernel_size, stride, padding))
        elif cv_type == "transpose":
            block.add_module(cv_name, nn.ConvTranspose2d(in_channels,
out_channels, kernel_size, stride, padding, output_padding))
    if norm_config != None:
        norm_name, norm_type, num_features = norm_config
        if norm_type == "batch":
            block.add_module(norm_name, nn.BatchNorm2d(num_features))
        elif norm_type == "instance":
            block.add_module(norm_name,
nn.InstanceNorm2d(num_features))
    if dropout_config != None:
        dropout_name, dropout_p = dropout_config
        block.add_module(dropout_name, nn.Dropout2d(p=dropout_p))
    if relu_config != None:
        relu_name, relu_type, negative_slope = relu_config
        if relu_type == "relu":
            block.add_module(relu_name, nn.ReLU(inplace=True))
        elif relu_type == "leaky_relu":
            block.add_module(relu_name, nn.LeakyReLU(negative_slope,
inplace=True))
    if sigmoid_config != None:
        sigmoid_name = sigmoid_config
        block.add_module(sigmoid_name, nn.Sigmoid())
    return block

class Discriminator(nn.Module):
##TODO1 implement the discriminator (critic)
    def __init__(self):
        super().__init__()
        self.convBlock1 = convBlock(conv_config=("conv2d_1",
"default", 1, 128, 5, 2, 2, -1),
                           relu_config=("leaky_relu_1",
"leaky_relu", 0.1))
        self.convBlock2 = convBlock(conv_config=("conv2d_2",
"default", 128, 256, 5, 2, 2, -1),
                           relu_config=("leaky_relu_2",
"leaky_relu", 0.1))

```

```

        relu_config="leaky_relu_2",
"leaky_relu", 0.1))
        self.convBlock3 = convBlock(conv_config="conv2d_3",
"default", 256, 512, 5, 2, 2, -1),
                           relu_config="leaky_relu_2",
"leaky_relu", 0.1))
        self.linear1 = nn.Linear(in_features=512*4*4, out_features=1)
def forward(self, x):
    output = self.convBlock1(x)
    output = self.convBlock2(output)
    output = self.convBlock3(output)
    output = torch.flatten(output, 1)
    output = self.linear1(output)
    return output

class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.linear1 = nn.Linear(in_features=128,
out_features=512*4*4)
        self.convBlock1 = convBlock(conv_config="conv2d_1",
"transpose", 512, 256, 5, 1, 0, 0),
                           norm_config="batchNorm_1",
"batch", 256),
                           relu_config="relu_1", "relu",
None))
        self.convBlock2 = convBlock(conv_config="conv2d_2",
"transpose", 256, 128, 5, 2, 2, 1),
                           norm_config="batchNorm_1",
"batch", 128),
                           relu_config="relu_1", "relu",
None))
        self.convBlock3 = convBlock(conv_config="conv2d_3",
"transpose", 128, 1, 5, 2, 2, 1),
                           sigmoid_config="sigmoid_1"))
    def forward(self, x):
        output = torch.reshape(self.linear1(x), (-1, 512, 4, 4))
        output = self.convBlock1(output)
        output = self.convBlock2(output)
        output = self.convBlock3(output)
        output = nn.functional.interpolate(output, scale_factor=28/32,
mode='bilinear', align_corners=True, recompute_scale_factor=True)
        return output
discriminator = Discriminator().cuda()
generator = Generator().cuda()

```

Network verification

TODO 3: What is the input and output shape of the generator and discriminator network? Verify that the implemented networks are the same as the answer you have provided.

```
print("Discriminator")
print("in: (1, 28, 28), out: (1, )")
summary(discriminator, input_size=trainX[0].shape)
print("Generator")
print("in: (128, ), out: (1, 28, 28)")
summary(generator, input_size=(128,))

Discriminator
in: (1, 28, 28), out: (1, )
=====
Layer (type:depth-idx)          Param #
=====
└ Sequential: 1-1
    └ Conv2d: 2-1           3,328
    └ LeakyReLU: 2-2         --
└ Sequential: 1-2
    └ Conv2d: 2-3           819,456
    └ LeakyReLU: 2-4         --
└ Sequential: 1-3
    └ Conv2d: 2-5           3,277,312
    └ LeakyReLU: 2-6         --
└ Linear: 1-4              8,193
=====
Total params: 4,108,289
Trainable params: 4,108,289
Non-trainable params: 0
=====
Generator
in: (128, ), out: (1, 28, 28)
=====
Layer (type:depth-idx)          Param #
=====
└ Linear: 1-1                 1,056,768
└ Sequential: 1-2
    └ ConvTranspose2d: 2-1   3,277,056
    └ BatchNorm2d: 2-2       512
    └ ReLU: 2-3              --
└ Sequential: 1-3
    └ ConvTranspose2d: 2-4   819,328
    └ BatchNorm2d: 2-5       256
    └ ReLU: 2-6              --
└ Sequential: 1-4
    └ ConvTranspose2d: 2-7   3,201
    └ Sigmoid: 2-8           --
```

```
=====
Total params: 5,157,121
Trainable params: 5,157,121
Non-trainable params: 0
=====

=====
Layer (type:depth-idx)           Param #
=====
└─Linear: 1-1                   1,056,768
└─Sequential: 1-2               --
   └─ConvTranspose2d: 2-1       3,277,056
   └─BatchNorm2d: 2-2          512
   └─ReLU: 2-3                 --
└─Sequential: 1-3               --
   └─ConvTranspose2d: 2-4       819,328
   └─BatchNorm2d: 2-5          256
   └─ReLU: 2-6                 --
└─Sequential: 1-4               --
   └─ConvTranspose2d: 2-7       3,201
   └─Sigmoid: 2-8              --
=====

Total params: 5,157,121
Trainable params: 5,157,121
Non-trainable params: 0
=====
```

Parameter Initialization

After the network is initialized, we then set up training hyperparameters for the training. In this part, hyperparameters have already been partially provided in the cell below, though some of them are intentionally left missing (`None`). Your task is to fill the missing parameters based on the pseudocode above.

TODO4: Initialize the missing model hyperparameters and optimizers based on the pseudocode above.

Note: To hasten the training process of our toy experiment, the training step and batch size is reduced to 3000 and 32, respectively.

```
import torch.optim as optim
NUM_ITERATION = 3000
BATCH_SIZE = 32
fixed_z = torch.randn((8, 128)).cuda()
def schedule(i):
    lr = 1e-4
    if(i > 2500): lr *= 0.1
    return lr
losses = {'D' : [], 'G' : []}
```

```

## TODO4 initialize missing hyperparameter and optimizer
G_optimizer = optim.Adam(generator.parameters(), lr=1e-4, betas=(0.0,
0.9)) # Assuming generator is defined elsewhere
D_optimizer = optim.Adam(discriminator.parameters(), lr=1e-4,
betas=(0.0, 0.9)) # Assuming discriminator is defined elsewhere
GP_lambda = 10 # Gradient penalty lambda
n_critic = 5 # Number of discriminator updates per generator update

```

Data preparation

TODO 5: Create a dataloader that could generate the data in line 4. The dataloader should return x, z, ϵ with a batch size of `BATCH_SIZE`

```

# TODO5 implement dataloader
from torch.utils.data import DataLoader, Dataset

class MNISTDataset(Dataset):
    def __init__(self, x):
        self.x = x.astype(np.float32)

    def __getitem__(self, index):
        x = self.x[index]
        z = np.random.uniform(low=0.0, high=1.0,
size=128).astype(np.float32)
        e = np.float32(np.random.uniform(low=0.0, high=1.0))
        return x, z, e

    def __len__(self):
        return self.x.shape[0]

train_dataset = MNISTDataset(trainX)
train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE,
shuffle=True, pin_memory=True)

```

Training loop

This section is the place where the training section starts. It is highly recommended that you understand the pseudocode before performing the tasks below. To train the WGAN-GP you have to perform the following tasks: TODO6: Update the learning rate base on the provided scheduler. TODO7: Sample the data from the dataloader (Line 4). TODO8 : Calcualte the discriminator loss (Line 5-7).

- In the line 7 you have to implement the gradient penalty term $\lambda \left(\left\| \nabla_{x^\wedge} D_w(x^\square \circ ^\wedge) \right\|_2 - 1 \right)^2$, which is a custom gradient. You may read <https://pytorch.org/docs/stable/generated/torch.autograd.grad.html> to find how custom gradient is implemented.
- HINT: Gradient norm calculation is still part of the computation graph.

TODO9: Update the discriminator loss (Line 9). TODO10: Calculate and update the generator loss (Line 11-12).

If your implementation is correct, the generated images should resemble an actual digit character after 500 iterations.

```
from tqdm import tqdm
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
for i in tqdm(range(NUM_ITERATION)):
    ## TODO6 update learning rate
    generator.zero_grad()
    D_optimizer.param_groups[0]["lr"] = schedule(i)
    G_optimizer.param_groups[0]["lr"] = schedule(i)

    for t in range(n_critic):
        ## TODO7 line 4: sample data from dataloader
        x, z, e = next(iter(train_loader))
        x = x.to(device, dtype=torch.float)
        z = z.to(device, dtype=torch.float)
        e = torch.reshape(e, (BATCH_SIZE, 1, 1, 1)).to(device,
        dtype=torch.float)

        ## TODO8 line5-7 : calculate discriminator loss
        x_tilde = generator(z)
        x_hat = e*x + (1-e)*x_tilde
        D_x = discriminator(x)
        D_x_hat = discriminator(x_hat)
        D_x_tilde = discriminator(x_tilde)

        discriminator.zero_grad()
        grad = torch.autograd.grad(
            outputs = D_x_hat,
            inputs = x_hat,
            grad_outputs = torch.ones(D_x_hat.size(), device = device),
            create_graph = True,
            retain_graph = True
        )[0]
        grad_penalty = GP_lambda*((grad.norm(2, dim=(2,3)) - 1)**2)
        D_loss = (D_x_tilde - D_x + grad_penalty).mean()
        losses['D'].append(D_loss.cpu().detach().numpy())

        ## TODO9 : line 9 update discriminator loss
        D_loss.backward()
        D_optimizer.step()
    ## TODO10 : line 11-12 calculate and update the generator loss
    generator.zero_grad()
    _, z, _ = next(iter(train_loader))
    z = z.to(device, dtype=torch.float)
```

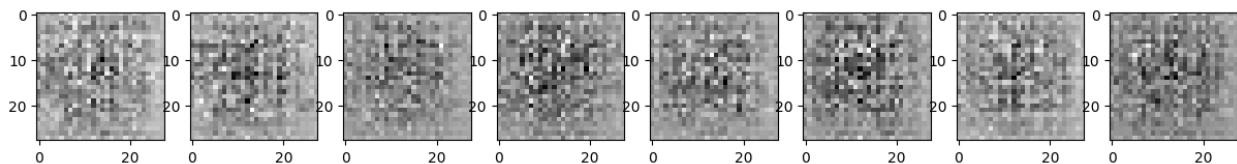
```

G_loss = (-discriminator(generator(z))).mean()
losses['G'].append(G_loss.cpu().detach().numpy())
G_loss.backward()
G_optimizer.step()

# Output visualization : If your reimplementation is correct, the
generated images should start resembling a digit character after 500
iterations.
if(i % 100 == 0):
    plt.figure(figsize = (15,75))
    print(losses['D'][-1], losses['G'][-1])
    with torch.no_grad():
        res = generator(fixed_z).cpu().detach().numpy()
    for k in range(8):
        plt.subplot( int('18{}'.format(k+1)) )
        plt.imshow( res[k].transpose(1, 2, 0)[..., 0], cmap = 'gray' )
    plt.show()

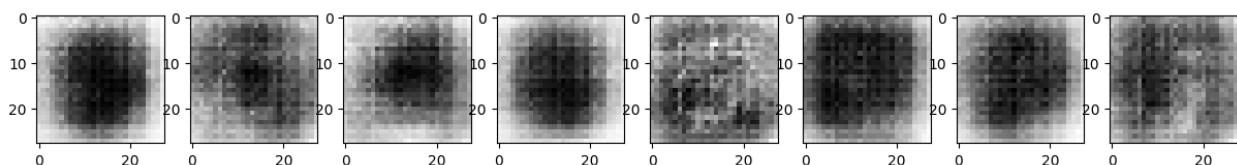
0% | 0/3000 [00:00<?, ?it/s]
5.8275185 1.940905

```



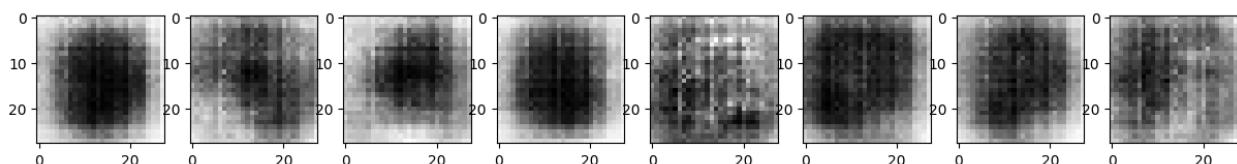
3% | 100/3000 [00:28<12:07, 3.98it/s]

-7.21551 -5.1610303



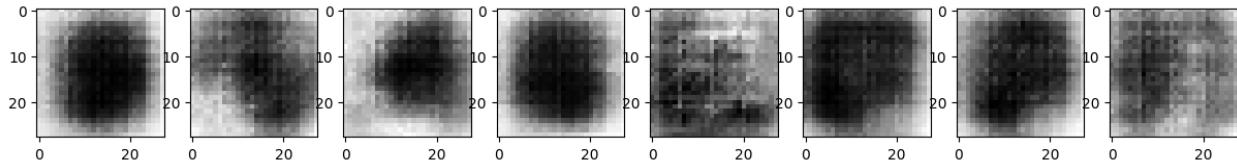
7% | 200/3000 [00:54<11:56, 3.91it/s]

-7.5414596 -3.5987718



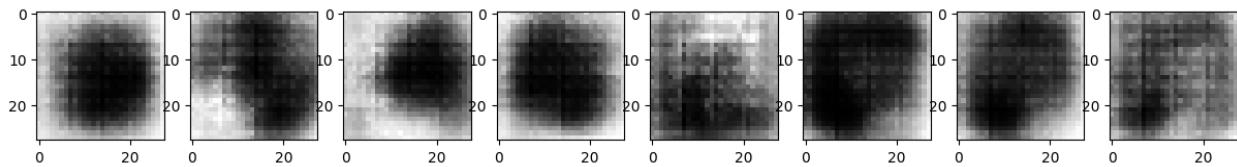
10% | 300/3000 [01:21<11:48, 3.81it/s]

-7.649962 -1.5924876



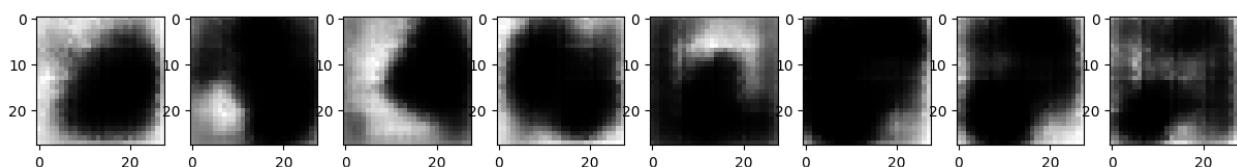
13% | | 400/3000 [01:48<11:39, 3.71it/s]

-8.024864 -2.2697873



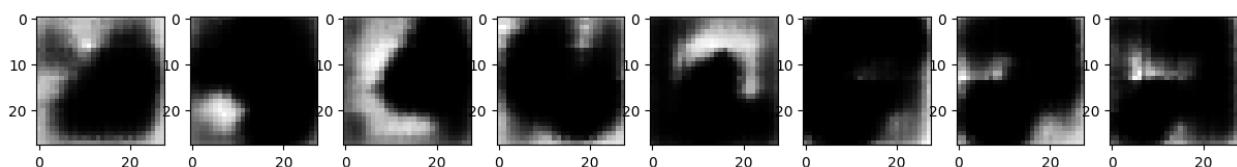
17% | | 500/3000 [02:16<11:24, 3.65it/s]

-5.864418 -1.246404



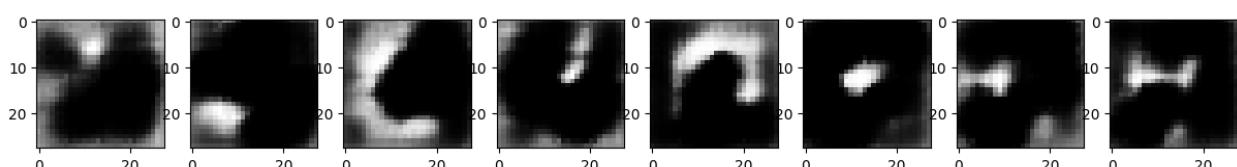
20% | | 600/3000 [02:44<10:51, 3.68it/s]

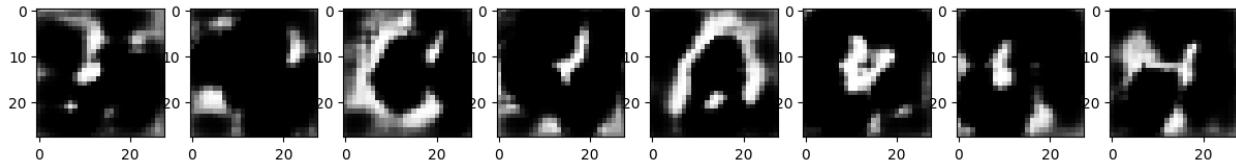
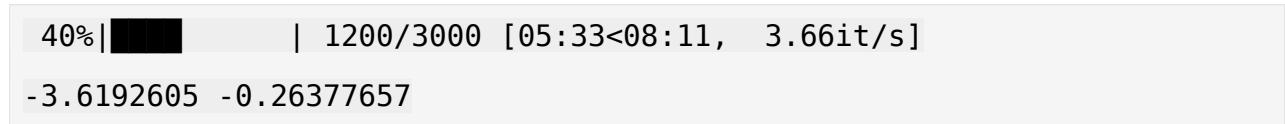
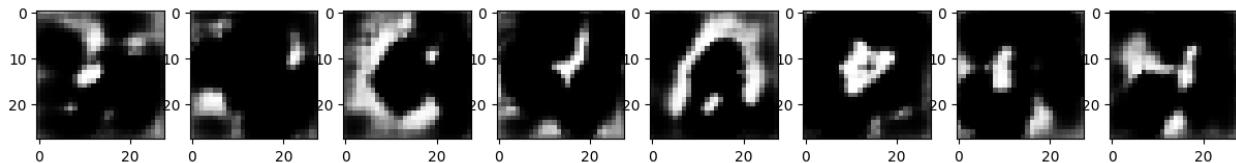
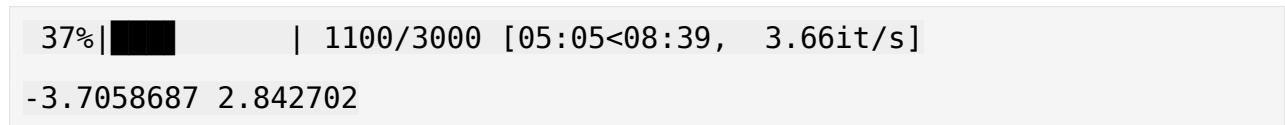
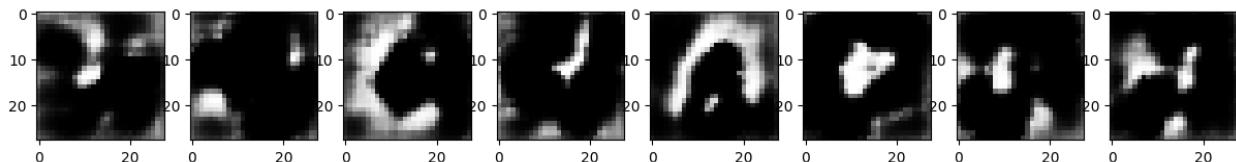
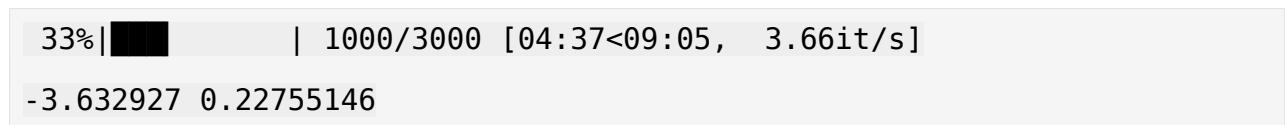
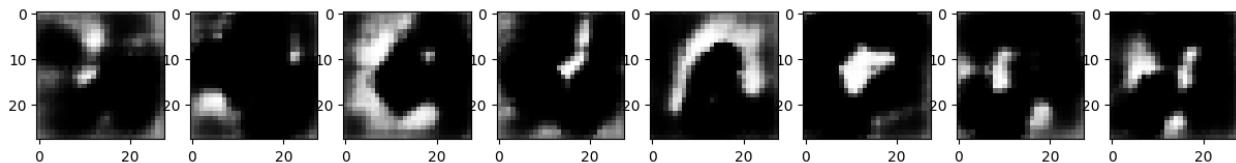
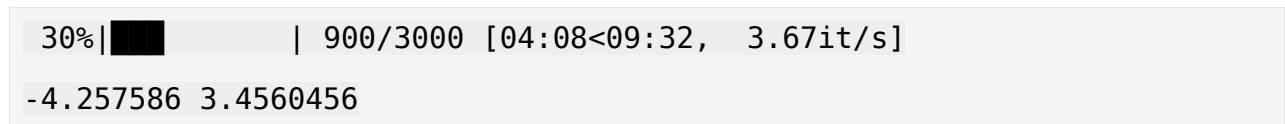
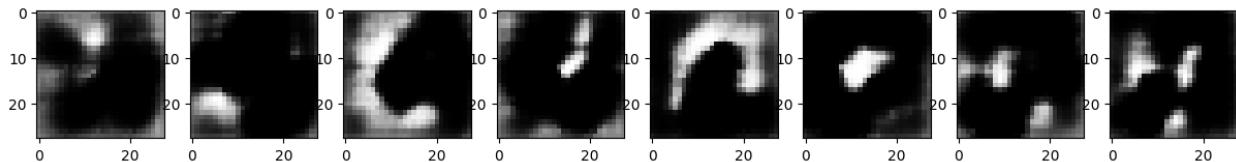
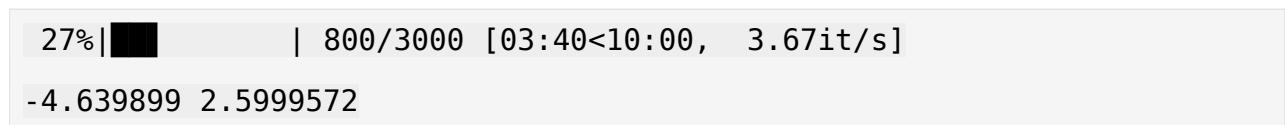
-6.1231833 2.2280574



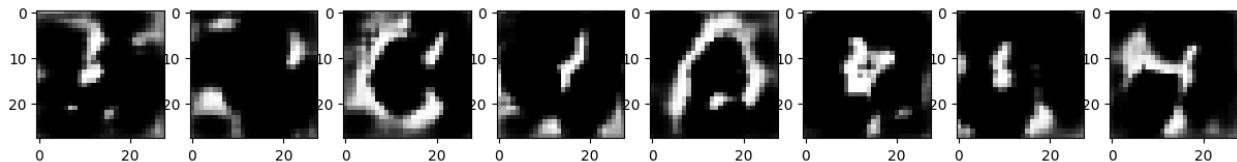
23% | | 700/3000 [03:12<10:22, 3.69it/s]

-6.0145864 2.404243

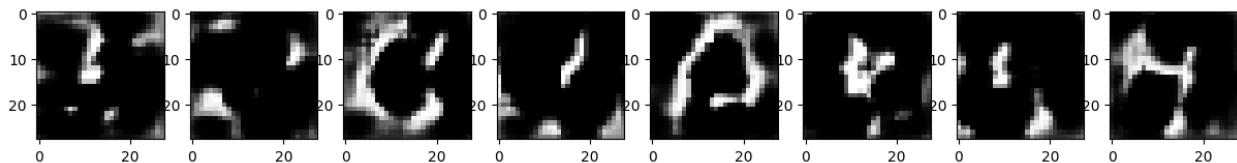




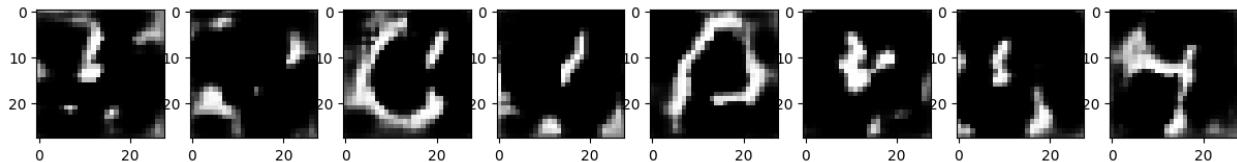
43% | [██████] | 1300/3000 [06:01<07:45, 3.66it/s]
-2.7046258 1.4193333



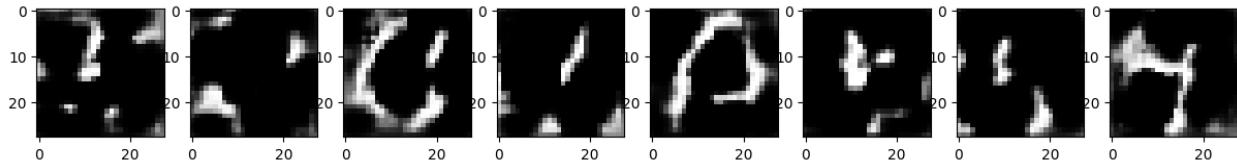
47% | [██████] | 1400/3000 [06:29<07:18, 3.65it/s]
-3.029561 1.6745133



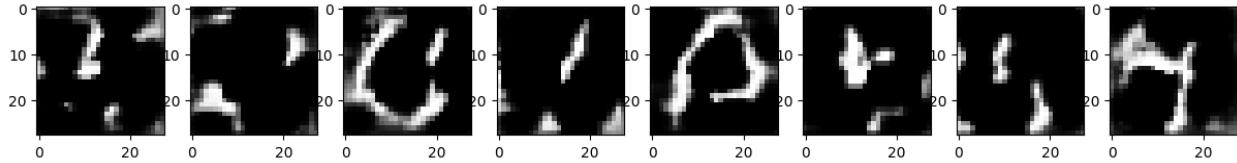
50% | [██████] | 1500/3000 [06:57<06:48, 3.67it/s]
-1.7415005 0.18732546



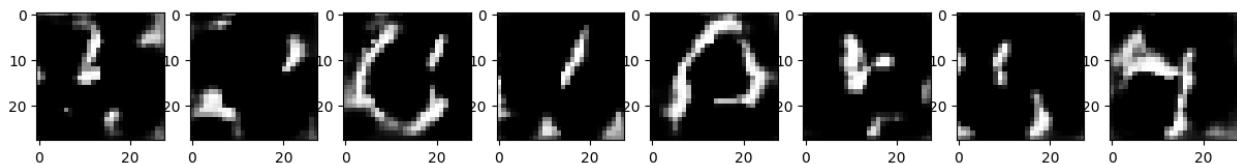
53% | [██████] | 1600/3000 [07:26<06:22, 3.66it/s]
-2.4372911 0.8195487



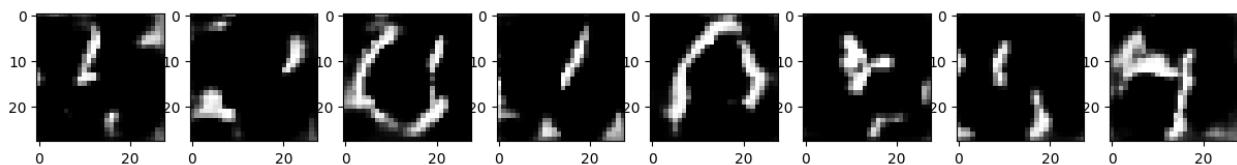
57% | [██████] | 1700/3000 [07:54<05:54, 3.66it/s]
-3.0422974 -1.1131147



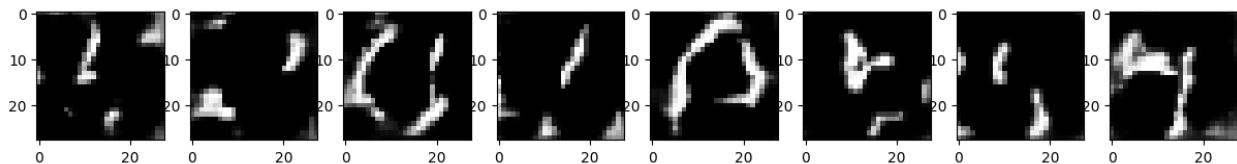
60% | [██████] | 1800/3000 [08:22<05:25, 3.68it/s]
-2.7214675 1.0587144



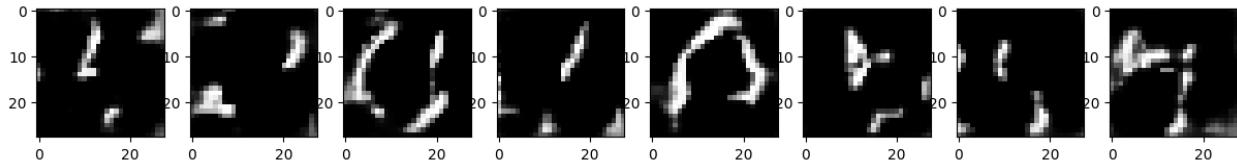
63% | [██████] | 1900/3000 [08:50<04:59, 3.68it/s]
-2.3935833 2.075819



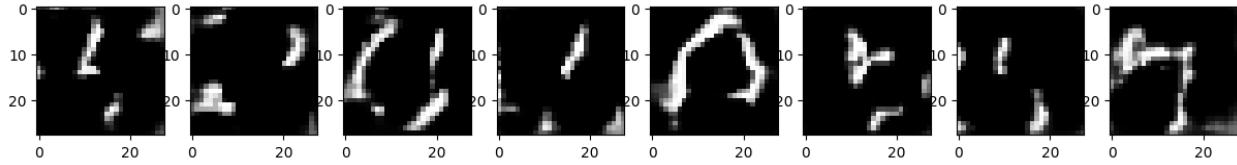
67% | [██████] | 2000/3000 [09:18<04:33, 3.66it/s]
-2.2758913 0.5534994



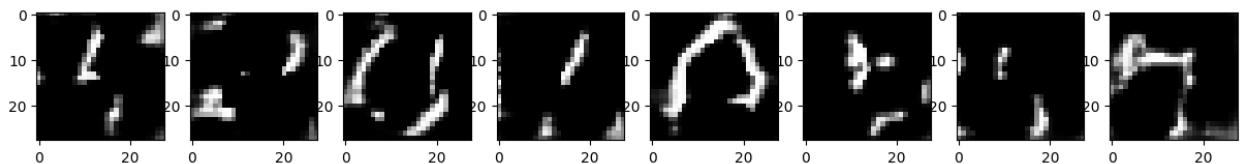
70% | [██████] | 2100/3000 [09:46<04:05, 3.66it/s]
-2.6100245 2.3330238



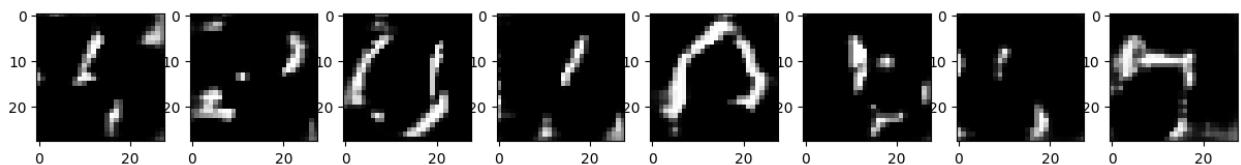
73% | [██████] | 2200/3000 [10:14<03:38, 3.66it/s]
-2.5094275 2.4516563



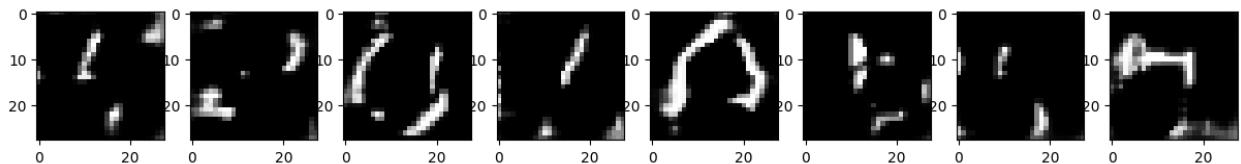
77% | [██████] | 2300/3000 [10:42<03:11, 3.66it/s]
-2.4488578 1.5001588



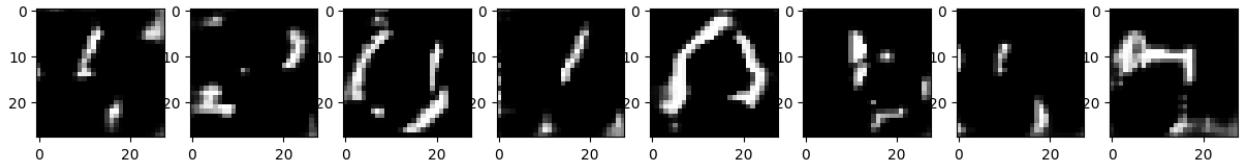
80% | [██████] | 2400/3000 [11:11<02:43, 3.67it/s]
-2.5207345 0.70771563



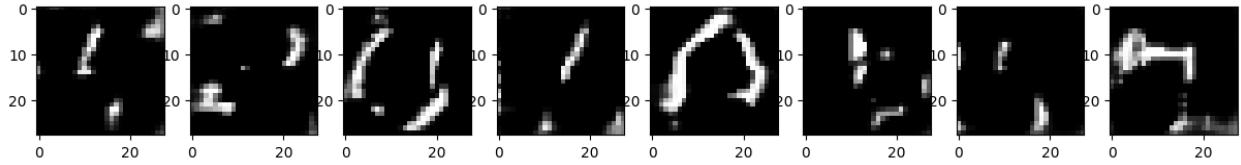
83% | [██████] | 2500/3000 [11:39<02:16, 3.67it/s]
-2.419629 0.2244496

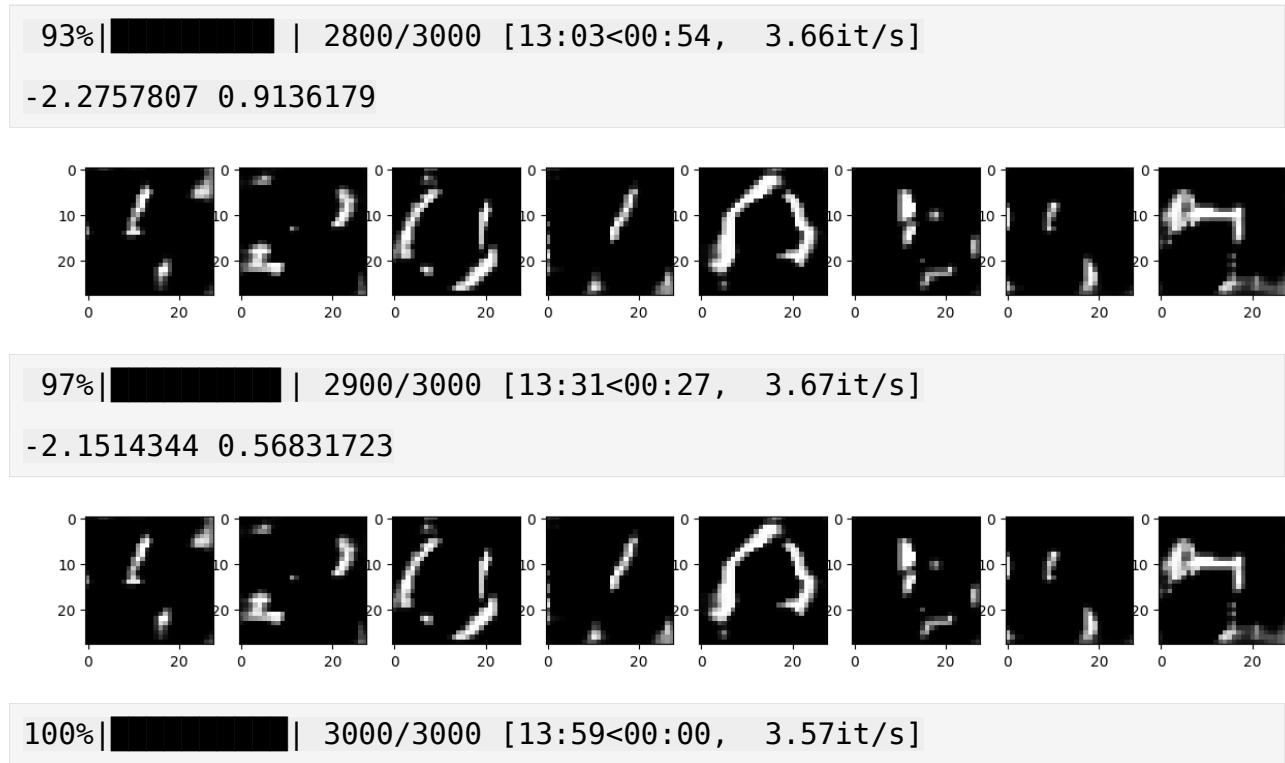


87% | [██████] | 2600/3000 [12:07<01:49, 3.66it/s]
-2.3123913 0.63404596



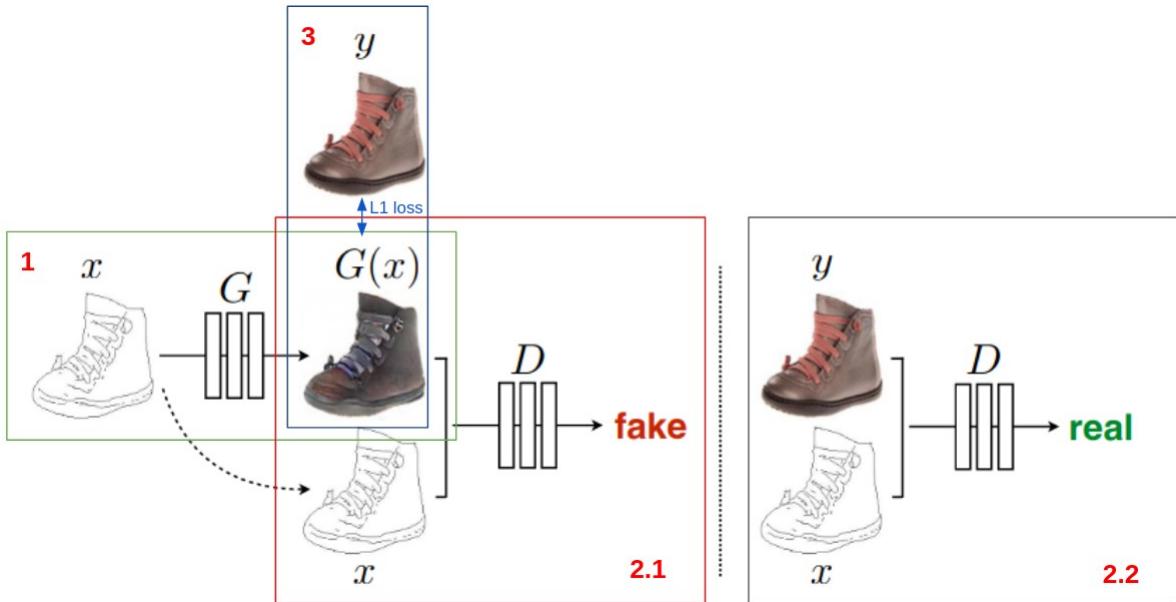
90% | [██████] | 2700/3000 [12:35<01:21, 3.66it/s]
-2.248591 0.4311853





Part 2 : pix2pix reimplementaion (cGAN on paired image translation)

In this exercise, we are reimplementing a paired image translation model, an application of a generative adversarial network (GAN). The model we are going to implement is pix2pix (<https://arxiv.org/pdf/1611.07004.pdf>), one of the earliest paired image translation models based on GAN. The pipeline of pix2pix is shown in the Figure below.



From the figure above, the pipeline consists of three main parts:

- a. Generation phase : the generator G create the generated image $G(x)$ from the given input x .
- a. Discrimination phase :

In step 2.1, the discriminator D receives an input image x and the generated image $G(x)$, then the discriminator has to learn to predict that the generated image $G(x)$ is fake. In step 2.2, the discriminator D receives an input image x and the ground truth image y , then the discriminator has to predict that the image y is real.
- a. Refinement phase: Refine the quality of the generated image $G(x)$ by encouraging the generated image to be close to an actual image y by using L1 as an objective.

The objective of pix2pix is to train an optimal generator G base on the objective function :

$$G = \arg \min_G \max_D L_{cGAN}(G, D) + \lambda L_1(G)$$

- The term $\arg \min_G \max_D L_{cGAN}(G, D)$ is the objective function of the first and second step, which is a standard cGAN loss : $L_{cGAN}(G, D) = E_{\{x,y\}}[\log D(x,y)] + E_{\{x,z\}}[\log(1 - D(x, G(x,z)))]$. The noise z is embedded in the generator in the form of dropout.
- The term $L_1(G)$ is the objective function of the third step where $L_1(G) = E_{x,y,z}[\|y - G(x, z)\|_1]$

The subsections will explain the dataset and training setup of this exercise.

Get dataset

```
!wget
http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar.gz
!tar -xzf facades.tar.gz

--2024-04-04 19:58:26--
http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar.gz
Resolving efrosgans.eecs.berkeley.edu (efrosgans.eecs.berkeley.edu)...
128.32.244.190
Connecting to efrosgans.eecs.berkeley.edu
(efrosgans.eecs.berkeley.edu)|128.32.244.190|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 30168306 (29M) [application/x-gzip]
Saving to: 'facades.tar.gz'

facades.tar.gz      100%[=====] 28.77M  1.70MB/s  in
15s

2024-04-04 19:58:46 (1.86 MB/s) - 'facades.tar.gz' saved
[30168306/30168306]
```

Import library

```
import cv2
import glob
import numpy as np
import torch
import torch.nn.functional as F
from torch import nn
from torchvision import transforms
import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec
```

Setting up facade dataset

The dataset chosen for this exercise is the CMP Facade Database which is a pair of facade images and its segmented component stored in RGB value. The objective of this exercise is to generate a facade given its simplified segmented component. Both input and output is a 256 x 256 RGB image normalized to [-1, 1].

```
train = (np.array([cv2.imread(i) for i in
glob.glob('facades/train/*')]), dtype = np.float32) /
255).transpose((0, 3, 1, 2))
train = (train - 0.5) * 2 #shift from [0,1] to [-1, 1]
trainX = train[:, :, :, :256]
trainY = train[:, :, :, :256]
```

```

val = (np.array([cv2.imread(i) for i in glob.glob('facades/val/*')]),
       dtype = np.float32) / 255).transpose(0, 3, 1, 2)
val = (val - 0.5) * 2 #shift from [0,1] to [-1, 1]
valX = val[:, :, :, 256:]
valY = val[:, :, :, :256]

print("Input size : {}, Output size = {}".format(trainX.shape,
trainY.shape))

Input size : (400, 3, 256, 256), Output size = (400, 3, 256, 256)

```

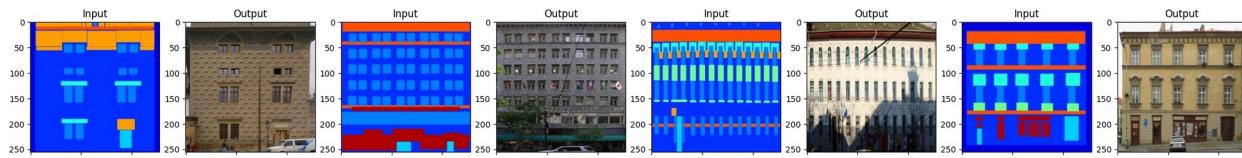
Dataset Visualization

```

import matplotlib.pyplot as plt
import numpy as np
plt.figure(figsize = (30,90))
for i in range(4):
    idx = np.random.randint(len(trainX))

    plt.subplot( int('19{}'.format(2*i+1)) )
    plt.title('Input')
    plt.imshow( (0.5 * trainX[idx].transpose((1, 2, 0)) + 0.5)[..., ::-1], cmap = 'gray' )
    plt.subplot( int('19{}'.format(2*i+2)) )
    plt.title('Output')
    plt.imshow( (0.5 * trainY[idx].transpose((1, 2, 0)) + 0.5)[..., ::-1], cmap = 'gray' )
plt.show()

```



Note

If you have trouble understanding the instruction provided in this homework or have any ambiguity about the instruction, you could also read the appendix section (section 6.1-6.2) in the paper for a detailed explanation.

Discriminator network

In this section, we are going to implement a discriminator network of pix2pix. The description of the discriminator network is provided in the Figure below.

Input size

256x256

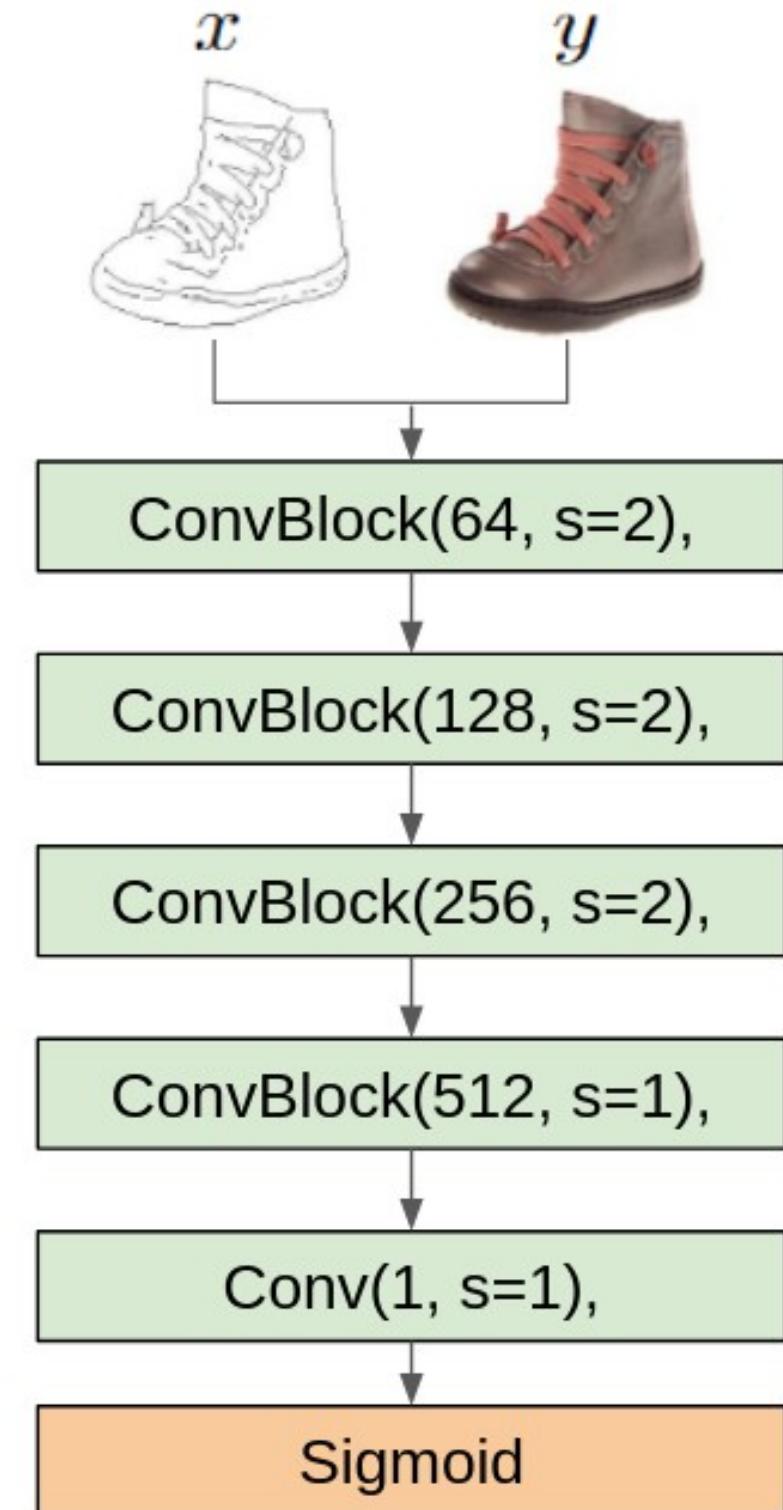
128x128

64x64

32x32

31x31

30x30



The network also has the following specific requirements:

- All convolutions are 4×4 spatial filters
- ConvBlock is a Convolution-InstanceNorm-ReLU layer
- InstanceNorm is not applied to the first C64 layer

- All ReLUs are leaky, with a slope of 0.2

TODO 11: Implement the discriminator network based on the description above.

TODO 12: What should be the size of the input and output of the discriminator for this task? Verify that the input and output of the implemented network are the same as the answer you have provided.

```
from torchsummary import summary
def convBlock(conv_config=None, norm_config=None, relu_config=None,
dropout_config=None, sigmoid_config=None):
    block = nn.Sequential()
    if conv_config != None:
        cv_name, cv_type, in_channels, out_channels, kernel_size,
        stride, padding, output_padding = conv_config
        if cv_type == "default":
            block.add_module(cv_name, nn.Conv2d(in_channels,
out_channels, kernel_size, stride, padding))
        elif cv_type == "transpose":
            block.add_module(cv_name, nn.ConvTranspose2d(in_channels,
out_channels, kernel_size, stride, padding, output_padding))
    if norm_config != None:
        norm_name, norm_type, num_features = norm_config
        if norm_type == "batch":
            block.add_module(norm_name, nn.BatchNorm2d(num_features))
        elif norm_type == "instance":
            block.add_module(norm_name,
nn.InstanceNorm2d(num_features))
    if dropout_config != None:
        dropout_name, dropout_p = dropout_config
        block.add_module(dropout_name, nn.Dropout2d(p=dropout_p))
    if relu_config != None:
        relu_name, relu_type, negative_slope = relu_config
        if relu_type == "relu":
            block.add_module(relu_name, nn.ReLU(inplace=False))
        elif relu_type == "leaky_relu":
            block.add_module(relu_name, nn.LeakyReLU(negative_slope,
inplace=False))
    if sigmoid_config != None:
        sigmoid_name = sigmoid_config
        block.add_module(sigmoid_name, nn.Sigmoid())
    return block

class Discriminator(nn.Module):
#TOD011 implement the discriminator network
    def __init__(self):
        super().__init__()
        self.convBlock1 = convBlock(conv_config=("conv2d_1",
"default", 6, 64, 4, 2, 1, -1),
        relu config=("leaky relu 1",
```

```

"leaky_relu", 0.2))
        self.convBlock2 = convBlock(conv_config="conv2d_2",
"default", 64, 128, 4, 2, 1, -1),
                                norm_config="instanceNorm_1",
"instance", 128),
                                relu_config="leaky_relu_2",
"leaky_relu", 0.2))
        self.convBlock3 = convBlock(conv_config="conv2d_3",
"default", 128, 256, 4, 2, 1, -1),
                                norm_config="instanceNorm_1",
"instance", 256),
                                relu_config="leaky_relu_3",
"leaky_relu", 0.2))
        self.convBlock4 = convBlock(conv_config="conv2d_4",
"default", 256, 512, 4, 1, 1, -1),
                                norm_config="instanceNorm_1",
"instance", 512),
                                relu_config="leaky_relu_4",
"leaky_relu", 0.2))
        self.conv = nn.Conv2d(in_channels=512, out_channels=1,
kernel_size=4, stride=1, padding=1)
        self.sigmoid = nn.Sigmoid()
    def forward(self, x):
        output = self.convBlock1(x)
        output = self.convBlock2(output)
        output = self.convBlock3(output)
        output = self.convBlock4(output)
        output = self.conv(output)
        output = self.sigmoid(output)
        return output
discriminator = Discriminator().cuda()

#TOD012 verify the discriminator
summary(discriminator, input_size = (6, 256, 256))

```

Layer (type:depth-idx)	Param #
Sequential: 1-1	--
└Conv2d: 2-1	6,208
└LeakyReLU: 2-2	--
Sequential: 1-2	--
└Conv2d: 2-3	131,200
└InstanceNorm2d: 2-4	--
└LeakyReLU: 2-5	--
Sequential: 1-3	--
└Conv2d: 2-6	524,544
└InstanceNorm2d: 2-7	--
└LeakyReLU: 2-8	--
Sequential: 1-4	--

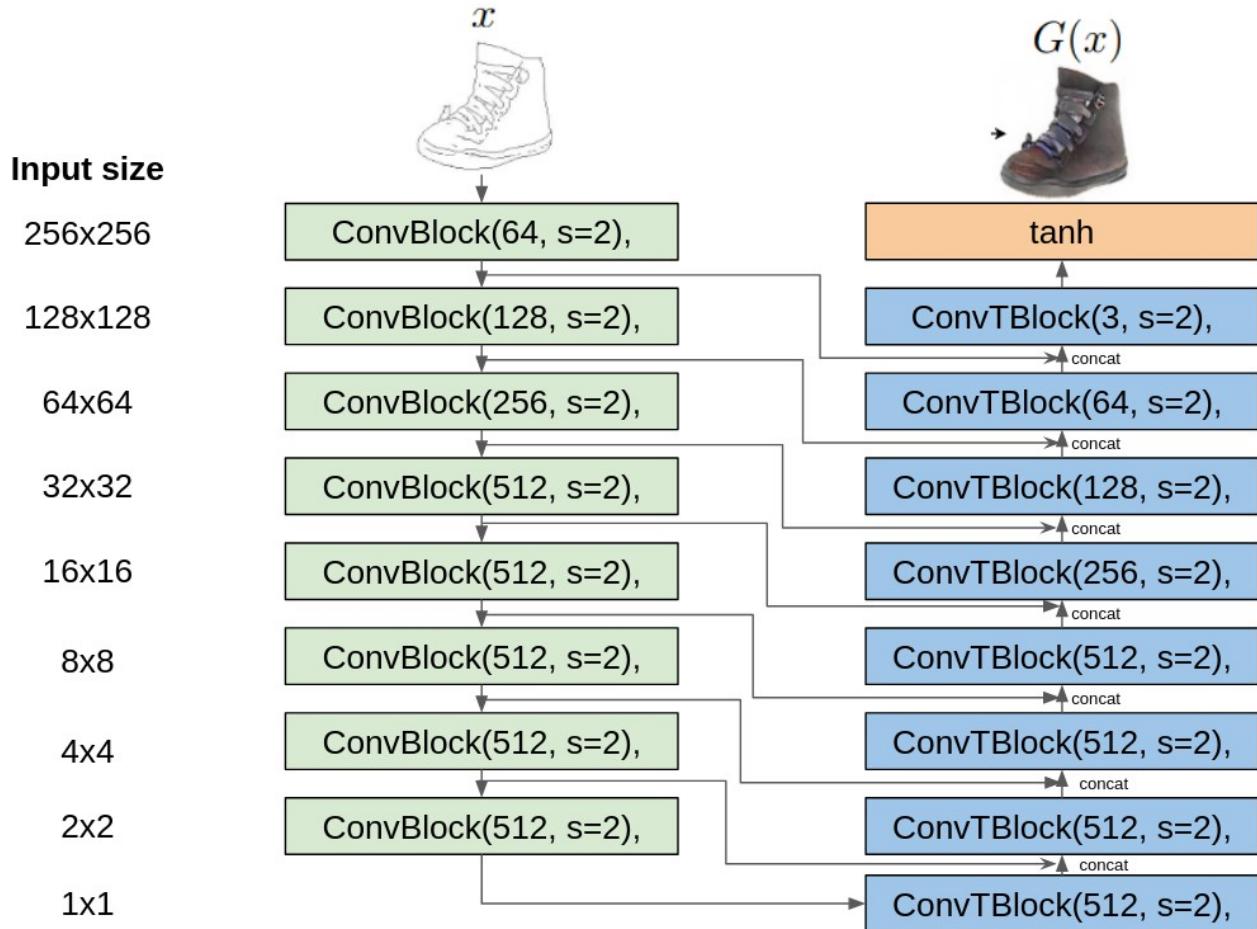
```

    └─Conv2d: 2-9                      2,097,664
    └─InstanceNorm2d: 2-10                --
    └─LeakyReLU: 2-11                  --
  ├─Conv2d: 1-5                      8,193
  └─Sigmoid: 1-6                     --
=====
Total params: 2,767,809
Trainable params: 2,767,809
Non-trainable params: 0
=====

=====
Layer (type:depth-idx)           Param #
=====
Sequential: 1-1                   --
  └─Conv2d: 2-1                  6,208
  └─LeakyReLU: 2-2                --
Sequential: 1-2                   --
  └─Conv2d: 2-3                  131,200
  └─InstanceNorm2d: 2-4                --
  └─LeakyReLU: 2-5                --
Sequential: 1-3                   --
  └─Conv2d: 2-6                  524,544
  └─InstanceNorm2d: 2-7                --
  └─LeakyReLU: 2-8                --
Sequential: 1-4                   --
  └─Conv2d: 2-9                  2,097,664
  └─InstanceNorm2d: 2-10                --
  └─LeakyReLU: 2-11                --
Conv2d: 1-5                      8,193
Sigmoid: 1-6                     --
=====
Total params: 2,767,809
Trainable params: 2,767,809
Non-trainable params: 0
=====
```

Generator network

In this section, we are going to implement a generator network of pix2pix. The generator is based on the U-NET based architecture (<https://arxiv.org/abs/1505.04597>). The Description of the generator network is provided in the Figure below.



The network also has the following specific requirements:

- All convolutions are 4×4 spatial filters
- ConvBlock is a Convolution-InstanceNorm-ReLU layer
- ConvTBlock is a ConvolutionTranspose-InstanceNorm-DropOut-ReLU layer with a dropout rate of 50%
- InstanceNorm is not applied to the first C64 layer in the encoder
- All ReLUs in the encoder are leaky, with a slope of 0.2, while ReLUs in the decoder are not leaky

TODO 13: Implement the generator network based on the description above.

TODO 14: What should be the size of the input and output of the generator for this task? Verify that the input and output of the implemented network are the same as the answer you have provided.

```
#HINT : you could also put multiple layers in a single list using
nn.ModuleList
class Example(nn.Module):
    def __init__(self):
        super().__init__()
        self.convs = nn.ModuleList([nn.Conv2d(25, 25, 3) for i in
```

```

range(5)])
def forward(self, x):
    for i in range(len(self.convs)):
        x = self.convs[i](x)
    return x
ex = Example().cuda()
print(ex(torch.zeros((8, 25, 32, 32)).cuda() ).shape
torch.Size([8, 25, 22, 22])

class Generator(nn.Module):
    #TODO13 implement the generator network
    def __init__(self):
        super().__init__()
        self.convBlocks =
nn.ModuleList([convBlock(conv_config=(f"conv2d_{i+5}", "default", 512,
512, 4, 2, 1, -1),
norm_config=(f"instanceNorm_{i+5}", "instance", 64),
relu_config=(f"leaky_relu_{i+5}", "leaky_relu", 0.2)) for i in
range(8)])
        self.convBlock1 = convBlock(conv_config=("conv2d_1", "default",
3, 64, 4, 2, 1, -1),
                                relu_config="leaky_relu_1",
"leaky_relu", 0.2))
        self.convBlock2 = convBlock(conv_config=("conv2d_2", "default",
64, 128, 4, 2, 1, -1),
                                norm_config="instanceNorm_2",
"instance", 128),
                                relu_config="leaky_relu_2",
"leaky_relu", 0.2))
        self.convBlock3 = convBlock(conv_config=("conv2d_3", "default",
128, 256, 4, 2, 1, -1),
                                norm_config="instanceNorm_3",
"instance", 256),
                                relu_config="leaky_relu_3",
"leaky_relu", 0.2))
        self.convBlock4 = convBlock(conv_config=("conv2d_4", "default",
256, 512, 4, 2, 1, -1),
                                norm_config="instanceNorm_4",
"instance", 512),
                                relu_config="leaky_relu_4",
"leaky_relu", 0.2))
        self.convBlocks =
nn.ModuleList([convBlock(conv_config=(f"conv2d_{i+5}", "default", 512,
512, 4, 2, 1, -1),
norm_config=(f"instanceNorm_{i+5}", "instance", 64),

```

```

relu_config=(f"leaky_relu_{i+5}", "leaky_relu", 0.2)) for i in
range(3)])
    self.convBlock5 = convBlock(conv_config="conv2d_8", "default",
512, 512, 4, 2, 1, -1),
                           relu_config=("leaky_relu_8",
"leaky_relu", 0.2))

    self.convTBlock1 = convBlock(conv_config="convT2d_1",
"transpose", 512, 512, 4, 2, 1, 0),
                           norm_config=("instanceNormD_1",
"instance", 512),
                           dropout_config=("dropout_1", 0.5),
                           relu_config=("relu_1", "relu", -1))

    self.convTBlocks =
nn.ModuleList([convBlock(conv_config=f"convT2d_{i+2}", "transpose",
512*2, 512, 4, 2, 1, 0),
norm_config=f"instanceNormD_{i+2}", "instance", 512),
dropout_config=f"dropout_{i+2}", 0.5),
relu_config=f"relu_{i+2}", "relu", -1) for i in range(3)])
    self.convTBlock2 = convBlock(conv_config="convT2d_5",
"transpose", 512*2, 256, 4, 2, 1, 0),
                           norm_config=("instanceNormD_5",
"instance", 256),
                           dropout_config=("dropout_5", 0.5),
                           relu_config=("relu_5", "relu", -1))

    self.convTBlock3 = convBlock(conv_config="convT2d_6",
"transpose", 256*2, 128, 4, 2, 1, 0),
                           norm_config=("instanceNormD_6",
"instance", 128),
                           dropout_config=("dropout_6", 0.5),
                           relu_config=("relu_6", "relu", -1))

    self.convTBlock4 = convBlock(conv_config="convT2d_7",
"transpose", 128*2, 64, 4, 2, 1, 0),
                           norm_config=("instanceNormD_7",
"instance", 64),
                           dropout_config=("dropout_7", 0.5),
                           relu_config=("relu_7", "relu", -1))

    self.convTBlock5 = convBlock(conv_config="convT2d_8",
"transpose", 64*2, 3, 4, 2, 1, 0),
                           norm_config=("instanceNormD_8",
"instance", 3))
    self.tanh = nn.Tanh()

def forward(self, x):
    output = self.convBlock1(x); cv1 = output
    output = self.convBlock2(output); cv2 = output
    output = self.convBlock3(output); cv3 = output
    output = self.convBlock4(output); cv4 = output

```

```

cv_s = []
for i in range(len(self.convBlocks)):
    output = self.convBlocks[i](output); cv_s.append(output)
output = self.convBlock5(output); cv5 = output

output = self.convTBlock1(output); cvT1 = output
cvT_s = []
for i in range(len(self.convTBlocks)):
    output = self.convTBlocks[i](torch.cat((output, cv_s[2-i]), dim=1)); cvT_s.append(output)
    output = self.convTBlock2(torch.cat((output, cv4), dim=1)); cvT2 = output
    output = self.convTBlock3(torch.cat((output, cv3), dim=1)); cvT3 = output
    output = self.convTBlock4(torch.cat((output, cv2), dim=1)); cvT4 = output
    output = self.convTBlock5(torch.cat((output, cv1), dim=1)); cvT5 = output
    output = self.tanh(output)
return output
generator = Generator().cuda()

#TOD014 verify the generator
summary(generator, input_size = (3, 256, 256))

```

Layer (type:depth-idx)	Param #
ModuleList: 1-1	--
└ Sequential: 2-1	--
└ Conv2d: 3-1	4,194,816
└ InstanceNorm2d: 3-2	--
└ LeakyReLU: 3-3	--
└ Sequential: 2-2	--
└ Conv2d: 3-4	4,194,816
└ InstanceNorm2d: 3-5	--
└ LeakyReLU: 3-6	--
└ Sequential: 2-3	--
└ Conv2d: 3-7	4,194,816
└ InstanceNorm2d: 3-8	--
└ LeakyReLU: 3-9	--
Sequential: 1-2	--
└ Conv2d: 2-4	3,136
└ LeakyReLU: 2-5	--
Sequential: 1-3	--
└ Conv2d: 2-6	131,200
└ InstanceNorm2d: 2-7	--
└ LeakyReLU: 2-8	--
Sequential: 1-4	--
└ Conv2d: 2-9	524,544

└ InstanceNorm2d: 2-10	--
└ LeakyReLU: 2-11	--
└ Sequential: 1-5	--
└ Conv2d: 2-12	2,097,664
└ InstanceNorm2d: 2-13	--
└ LeakyReLU: 2-14	--
└ Sequential: 1-6	--
└ Conv2d: 2-15	4,194,816
└ LeakyReLU: 2-16	--
└ Sequential: 1-7	--
└ ConvTranspose2d: 2-17	4,194,816
└ InstanceNorm2d: 2-18	--
└ Dropout2d: 2-19	--
└ ReLU: 2-20	--
└ ModuleList: 1-8	--
└ Sequential: 2-21	--
└ ConvTranspose2d: 3-10	8,389,120
└ InstanceNorm2d: 3-11	--
└ Dropout2d: 3-12	--
└ ReLU: 3-13	--
└ Sequential: 2-22	--
└ ConvTranspose2d: 3-14	8,389,120
└ InstanceNorm2d: 3-15	--
└ Dropout2d: 3-16	--
└ ReLU: 3-17	--
└ Sequential: 2-23	--
└ ConvTranspose2d: 3-18	8,389,120
└ InstanceNorm2d: 3-19	--
└ Dropout2d: 3-20	--
└ ReLU: 3-21	--
└ Sequential: 1-9	--
└ ConvTranspose2d: 2-24	4,194,560
└ InstanceNorm2d: 2-25	--
└ Dropout2d: 2-26	--
└ ReLU: 2-27	--
└ Sequential: 1-10	--
└ ConvTranspose2d: 2-28	1,048,704
└ InstanceNorm2d: 2-29	--
└ Dropout2d: 2-30	--
└ ReLU: 2-31	--
└ Sequential: 1-11	--
└ ConvTranspose2d: 2-32	262,208
└ InstanceNorm2d: 2-33	--
└ Dropout2d: 2-34	--
└ ReLU: 2-35	--
└ Sequential: 1-12	--
└ ConvTranspose2d: 2-36	6,147
└ InstanceNorm2d: 2-37	--
└ Tanh: 1-13	--

Layer (type:depth-idx)	Param #
ModuleList: 1-1	--
└ Sequential: 2-1	--
└ Conv2d: 3-1	4,194,816
└ InstanceNorm2d: 3-2	--
└ LeakyReLU: 3-3	--
└ Sequential: 2-2	--
└ Conv2d: 3-4	4,194,816
└ InstanceNorm2d: 3-5	--
└ LeakyReLU: 3-6	--
└ Sequential: 2-3	--
└ Conv2d: 3-7	4,194,816
└ InstanceNorm2d: 3-8	--
└ LeakyReLU: 3-9	--
Sequential: 1-2	--
└ Conv2d: 2-4	3,136
└ LeakyReLU: 2-5	--
Sequential: 1-3	--
└ Conv2d: 2-6	131,200
└ InstanceNorm2d: 2-7	--
└ LeakyReLU: 2-8	--
Sequential: 1-4	--
└ Conv2d: 2-9	524,544
└ InstanceNorm2d: 2-10	--
└ LeakyReLU: 2-11	--
Sequential: 1-5	--
└ Conv2d: 2-12	2,097,664
└ InstanceNorm2d: 2-13	--
└ LeakyReLU: 2-14	--
Sequential: 1-6	--
└ Conv2d: 2-15	4,194,816
└ LeakyReLU: 2-16	--
Sequential: 1-7	--
└ ConvTranspose2d: 2-17	4,194,816
└ InstanceNorm2d: 2-18	--
└ Dropout2d: 2-19	--
└ ReLU: 2-20	--
ModuleList: 1-8	--
└ Sequential: 2-21	--
└ ConvTranspose2d: 3-10	8,389,120
└ InstanceNorm2d: 3-11	--
└ Dropout2d: 3-12	--

└ReLU: 3-13	--
└Sequential: 2-22	--
└ConvTranspose2d: 3-14	8,389,120
└InstanceNorm2d: 3-15	--
└Dropout2d: 3-16	--
└ReLU: 3-17	--
└Sequential: 2-23	--
└ConvTranspose2d: 3-18	8,389,120
└InstanceNorm2d: 3-19	--
└Dropout2d: 3-20	--
└ReLU: 3-21	--
└Sequential: 1-9	--
└ConvTranspose2d: 2-24	4,194,560
└InstanceNorm2d: 2-25	--
└Dropout2d: 2-26	--
└ReLU: 2-27	--
└Sequential: 1-10	--
└ConvTranspose2d: 2-28	1,048,704
└InstanceNorm2d: 2-29	--
└Dropout2d: 2-30	--
└ReLU: 2-31	--
└Sequential: 1-11	--
└ConvTranspose2d: 2-32	262,208
└InstanceNorm2d: 2-33	--
└Dropout2d: 2-34	--
└ReLU: 2-35	--
└Sequential: 1-12	--
└ConvTranspose2d: 2-36	6,147
└InstanceNorm2d: 2-37	--
└Tanh: 1-13	--
<hr/>	
Total params:	54,409,603
Trainable params:	54,409,603
Non-trainable params:	0
<hr/>	

Data preparation

After the model is initialized, we then create a dataloader to sample the training data. In this paper, to sample the training data, you have to sequentially perform the following steps :

1. Randomly sample the data from the training set
2. Resizing both input and target to 286×286 .
3. Randomly cropping back both images to size 256×256 .
4. Random mirroring the images

TODO15: Implement a dataloader based on the description above. You are allowed to use the function in `torchvision.transforms` (<https://pytorch.org/vision/main/transforms.html>).

```

# TODO15 implement a dataloader
from torch.utils.data import DataLoader, Dataset

transform = transforms.Compose(
    [transforms.Resize([286, 286]),
     transforms.RandomCrop(256),
     transforms.RandomHorizontalFlip()])

class facadeDataset(Dataset):
    def __init__(self, x, y):
        self.x = x.astype(np.float32)
        self.y = y.astype(np.float32)

    def __getitem__(self, index):
        x = self.x[index] # Retrieve data
        y = self.y[index] # Retrieve data
        x = torch.tensor(x)
        y = torch.tensor(y)
        xy = torch.cat((x,y), dim=0)
        xy = transform(torch.tensor(xy))
        x, y = torch.split(xy, 3, dim=0)
        return x, y

    def __len__(self):
        return self.x.shape[0]

train_dataset = facadeDataset(trainX, trainY)
train_loader = DataLoader(train_dataset, batch_size=1, shuffle=True,
pin_memory=True)

```

Dataloader verification

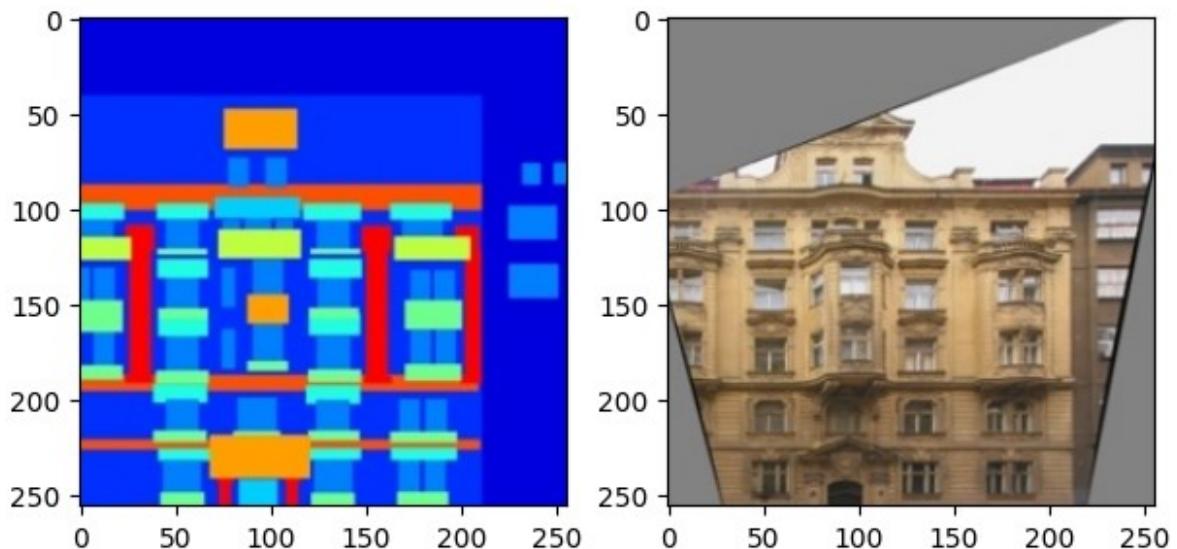
TODO16: Show that the implemented dataloader is working as intended. For instance, are both input and output are flipped and cropped correctly? To obtain a full score, you have to show at least eight data points.

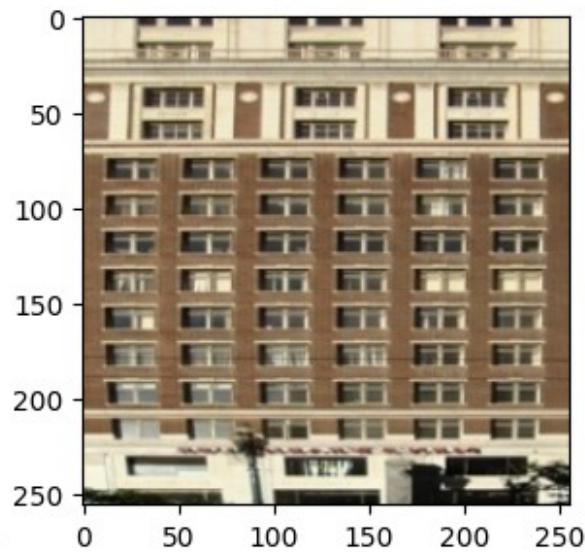
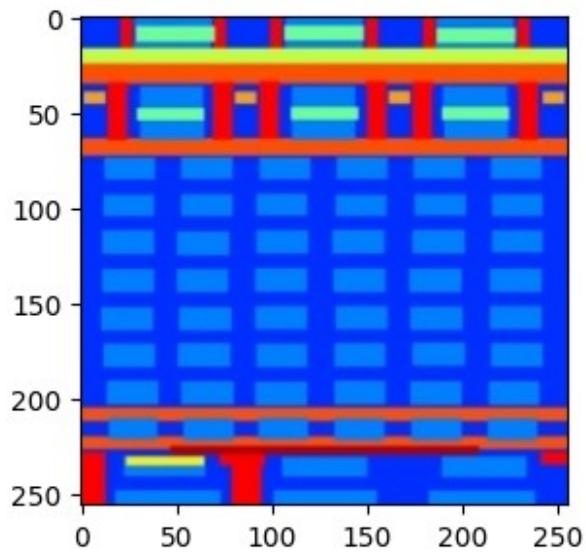
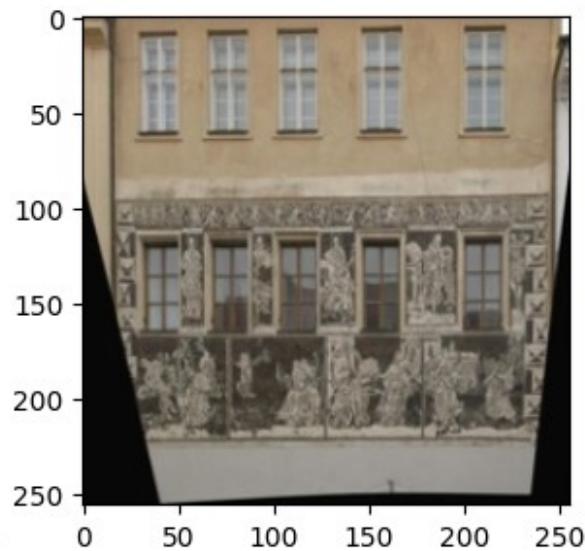
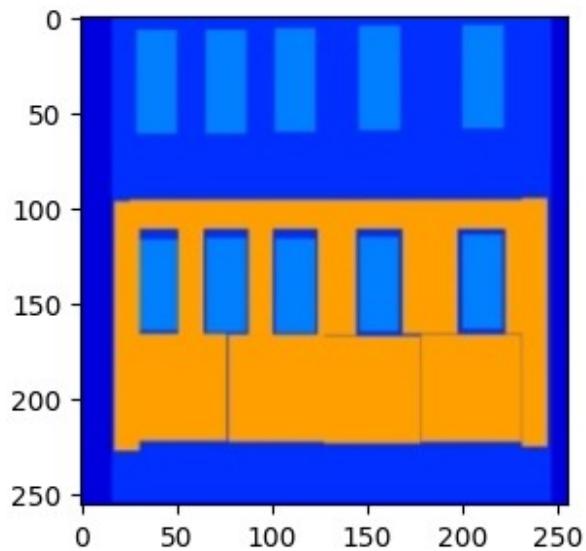
```

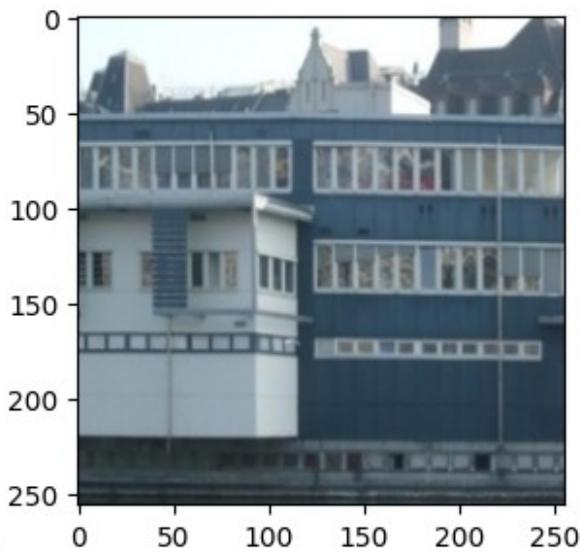
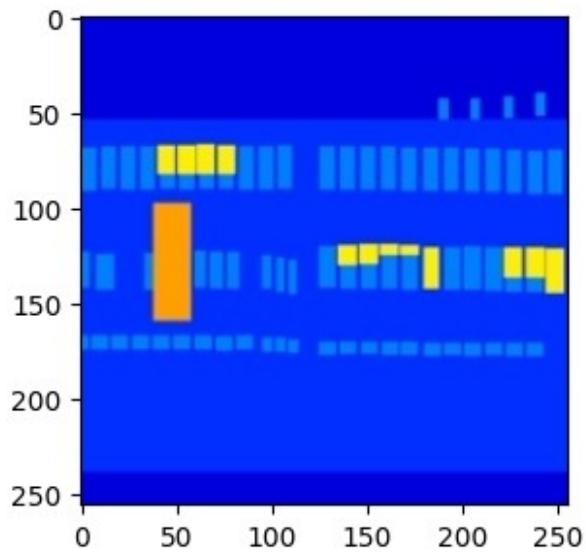
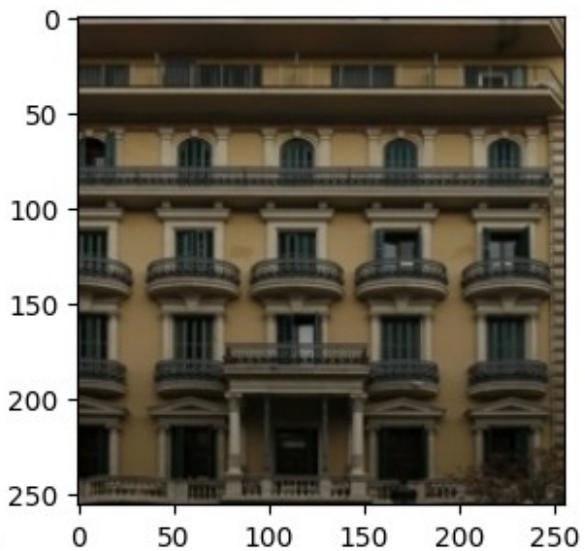
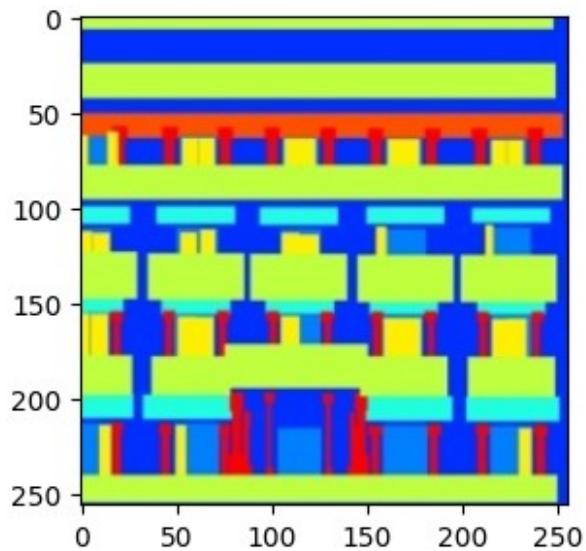
# TODO16 show that the dataloader is working properly
for i in range(8):
    plt.figure(figsize = (15,15))
    x, y = next(iter(train_loader))
    x = x[0]; y = y[0]
    plt.subplot(4, 4, 2*i+1)
    plt.imshow( (0.5 * x.numpy().transpose((1, 2, 0)) + 0.5)[..., ::-1] , cmap = 'gray' )
    plt.subplot(4, 4, 2*i+2)
    plt.imshow( (0.5 * y.numpy().transpose((1, 2, 0)) + 0.5)[..., ::-1] , cmap = 'gray' )
    plt.show()

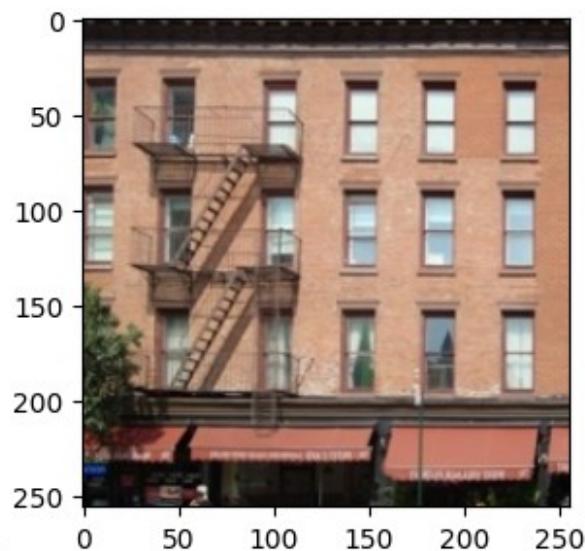
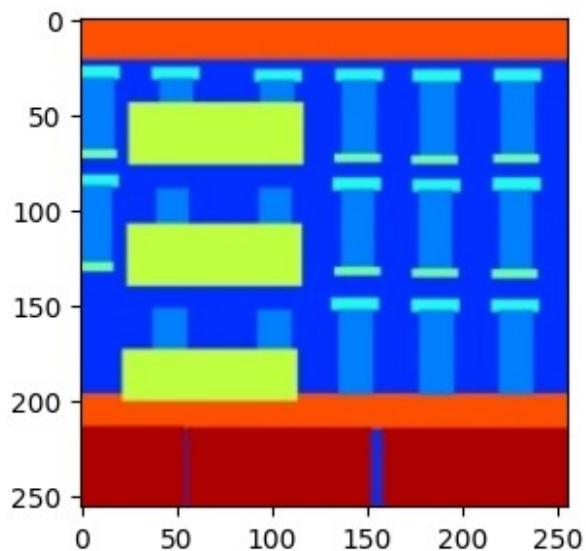
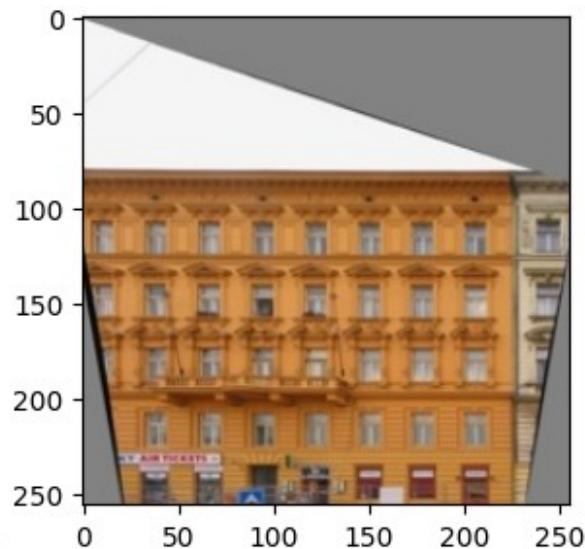
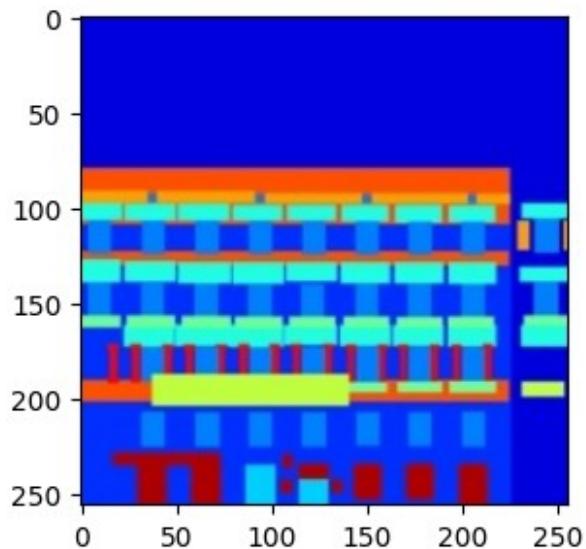
```

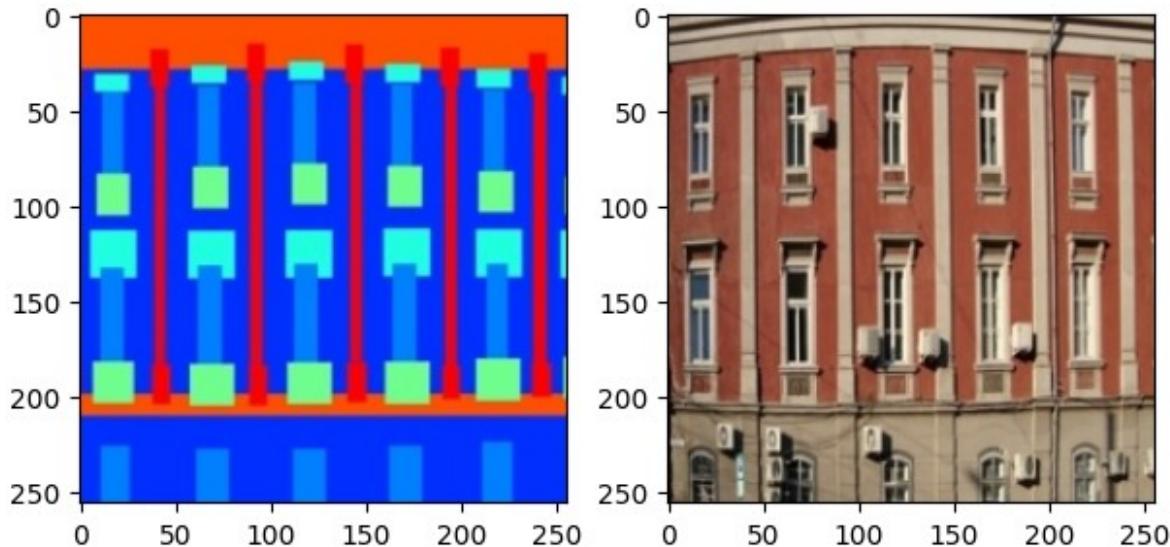
```
/tmp/ipykernel_34/46597281.py:20: UserWarning: To copy construct from
a tensor, it is recommended to use sourceTensor.clone().detach() or
sourceTensor.clone().detach().requires_grad_(True), rather than
torch.tensor(sourceTensor).
    xy = transform(torch.tensor(xy))
/opt/conda/lib/python3.10/site-packages/torchvision/transforms/function
nal.py:1603: UserWarning: The default value of the antialias parameter
of all the resizing transforms (Resize(), RandomResizedCrop(), etc.)
will change from None to True in v0.17, in order to be consistent
across the PIL and Tensor backends. To suppress this warning, directly
pass antialias=True (recommended, future default), antialias=None
(current default, which means False for Tensors and True for PIL), or
antialias=False (only works on Tensors - PIL will still use
antialiasing). This also applies if you are using the inference
transforms from the models weights: update the call to
weights.transforms(antialias=True).
    warnings.warn
```











Parameter Initialization

Model hyperparameters and optimizers have already been prepared.

```
import torch.optim as optim
from tqdm import tqdm
lr = 2e-4
LAMBDA = 100
BATCH_SIZE = 1
G_optimizer = optim.Adam(generator.parameters(), lr=lr, betas = (0.5, 0.999))
D_optimizer = optim.Adam(discriminator.parameters(), lr=lr, betas = (0.5, 0.999))
```

Training loop

The training process has the following specific requirements:

- The objective is divided by 2 while optimizing D , which slows down the rate at which D learns relative to G .
- This paper trains the generator G to maximize $\log D(x, G(x, z))$ instead of minimizing $\log(1 - D(x, G(x, z)))$ as the latter term does not provide sufficient gradient.

TODO17: Sample the data using the dataloader. TODO18: Calculate the discriminator loss and update the discriminator.

- During the update, the loss term $\log(1 - D(x, G(x, z)))$ contains both generator and discriminator. However, we only want to update the discriminator. Therefore, you have to detach the input from the generator graph first. Read <https://pytorch.org/docs/stable/generated/torch.Tensor.detach.html> for additional detail.

TODO19: Calculate the generator loss and update the generator.

HINT Hint 1: If you are struggling with this part, you could also read the PyTorch DCGAN tutorial as a guideline (https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html).

Hint 2: You could remove the L1 loss while debugging since the generator G could still generate the synthetic image even if the L1 loss is removed, though at the cost of increasing image artifacts.

Note The training schedule in this homework is only one eighth of the original schedule. It is expected that the generated image quality is worse than the one shown in the paper. Nevertheless, the generated facade should still resemble an actual one.

```
losses = {'D' : [None], 'G' : [None]}
loss_fn = nn.BCELoss()
loss_l1 = nn.L1Loss()
torch.autograd.set_detect_anomaly(True)
for i in tqdm(range(20001)):
    #TODO17 sample the data from the dataloader
    x, y = next(iter(train_loader))
    x = x.cuda()
    y = y.cuda()
    discriminator.zero_grad()
    #TODO18 calculate the discriminator loss and update the
    discriminator
    x_tilde = generator(x)
    D_x_gx = discriminator(torch.cat([x, x_tilde.detach()], dim=1))
    D_x_y = discriminator(torch.cat([x, y], dim=1))

    D_loss_zero = loss_fn(D_x_gx, torch.zeros(D_x_gx.shape).cuda())
    D_loss_one = loss_fn(D_x_y, torch.ones_like(D_x_y))
    D_loss = D_loss_zero + D_loss_one
    losses['D'].append(D_loss)
    D_loss.backward()
    D_optimizer.step()
    #TODO19 calculate the generator loss and update the generator
    generator.zero_grad()
    D_x_gx_tmp = discriminator(torch.cat([x, x_tilde], dim=1))
    G_loss = loss_fn(D_x_gx_tmp, torch.ones_like(D_x_gx_tmp)) +
    loss_l1(x_tilde, y)*LAMBDA
    losses['G'].append(G_loss)
    G_loss.backward()
    G_optimizer.step()
    # Output visualization : If your reimplementation is correct, the
    generated images should start resembling a facade after 2,500
    iterations
    if(i % 2500 == 0):
        with torch.no_grad():
            print(losses['D'][-1], losses['G'][-1])
            plt.figure(figsize = (40, 16))
```

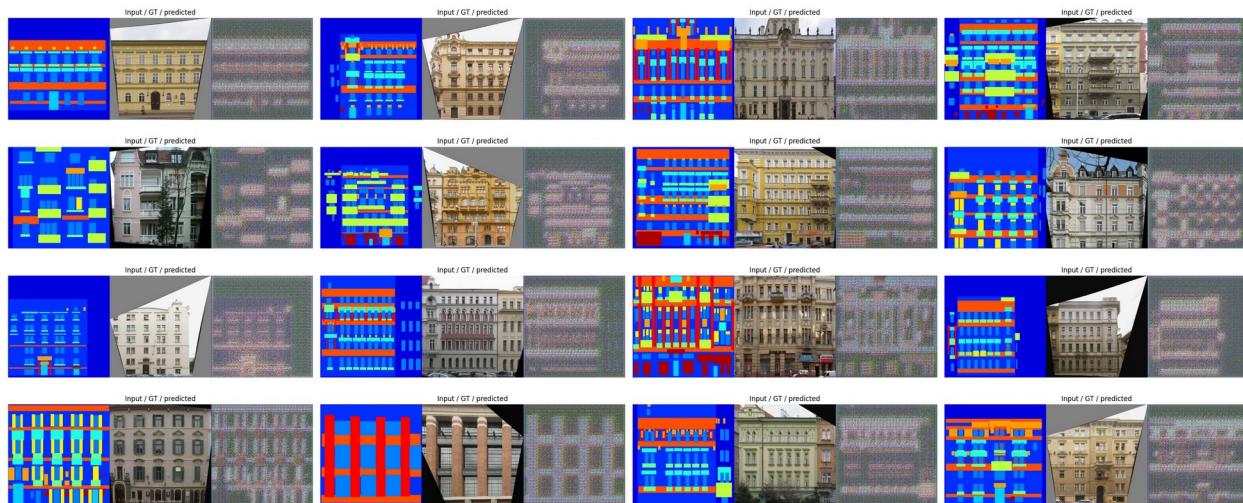
```

gs1 = gridspec.GridSpec(4, 4)
gs1.update(wspace=0.025)

sampleX_vis = 0.5 * valX[:16][:, ::-1, :, :] + 0.5
sampleY = 0.5 * valY[:16][:, ::-1, :, :] + 0.5
sampleX = torch.tensor(valX[:16]).cuda()
pred_val = 0.5 * generator(sampleX).cpu().detach().numpy()
[:, ::-1, :, :] + 0.5
vis = np.concatenate([sampleX_vis, sampleY, pred_val], axis = 3)
for i in range(vis.shape[0]):
    ax1 = plt.subplot(gs1[i])
    plt.title('Input / GT / predicted')
    plt.axis('off')
    plt.imshow( vis[i].transpose(1, 2, 0) )
plt.show()

0% | 0/20001 [00:00<?, ?it/s]/tmp/ipykernel_34/46597281.py:20: UserWarning: To copy construct
from a tensor, it is recommended to use sourceTensor.clone().detach()
or sourceTensor.clone().detach().requires_grad_(True), rather than
torch.tensor(sourceTensor).
    xy = transform(torch.tensor(xy))
/opt/conda/lib/python3.10/site-packages/torch/nn/modules/instancenorm.
py:80: UserWarning: input's size at dim=1 does not match num_features.
You can silence this warning by not passing in num_features, which is
not used because affine=False
    warnings.warn(f"input's size at dim={feature_dim} does not match
num_features. "
tensor(1.4235, device='cuda:0', grad_fn=<AddBackward0>)
tensor(62.2610, device='cuda:0', grad_fn=<AddBackward0>)

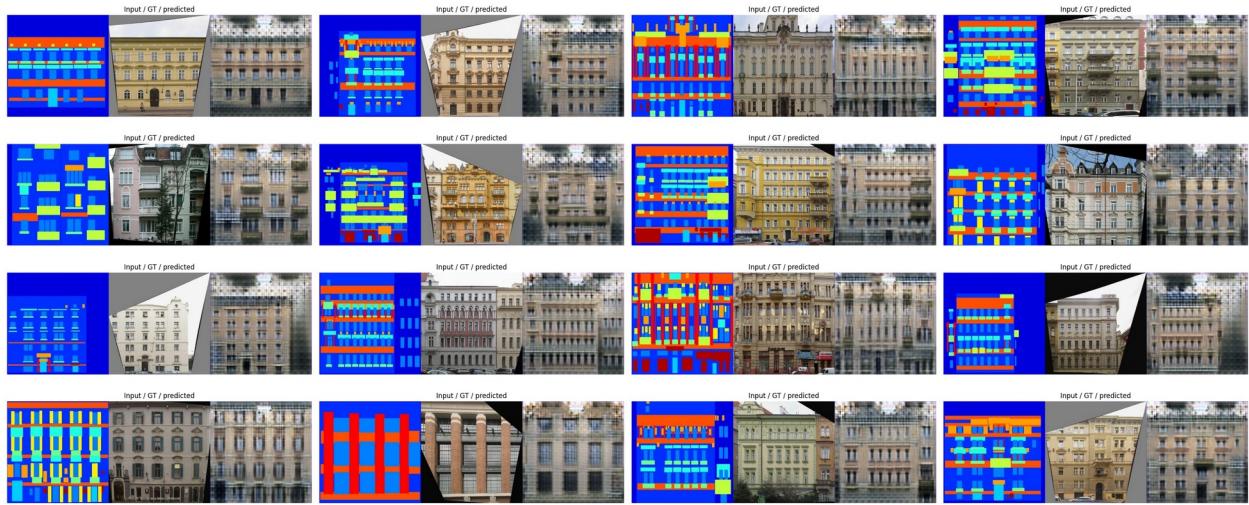
```



12% |

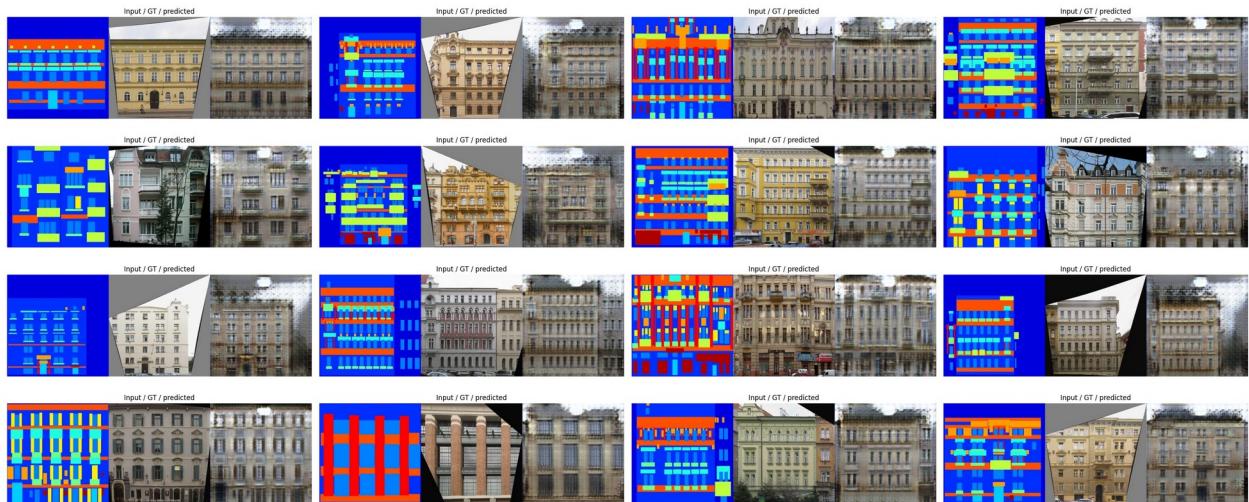
| 2500/20001 [06:07<42:27, 6.87it/s]

```
tensor(0.2607, device='cuda:0', grad_fn=<AddBackward0>)
tensor(45.3708, device='cuda:0', grad_fn=<AddBackward0>)
```



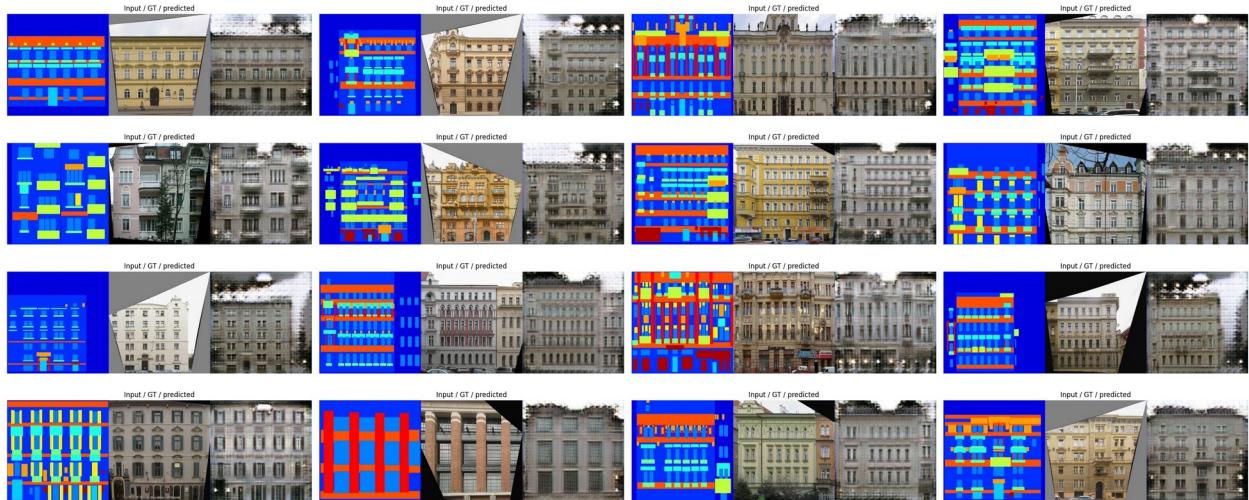
25% |  | 5000/20001 [12:14<35:41, 7.00it/s]

```
tensor(0.0310, device='cuda:0', grad_fn=<AddBackward0>)
tensor(35.0348, device='cuda:0', grad_fn=<AddBackward0>)
```



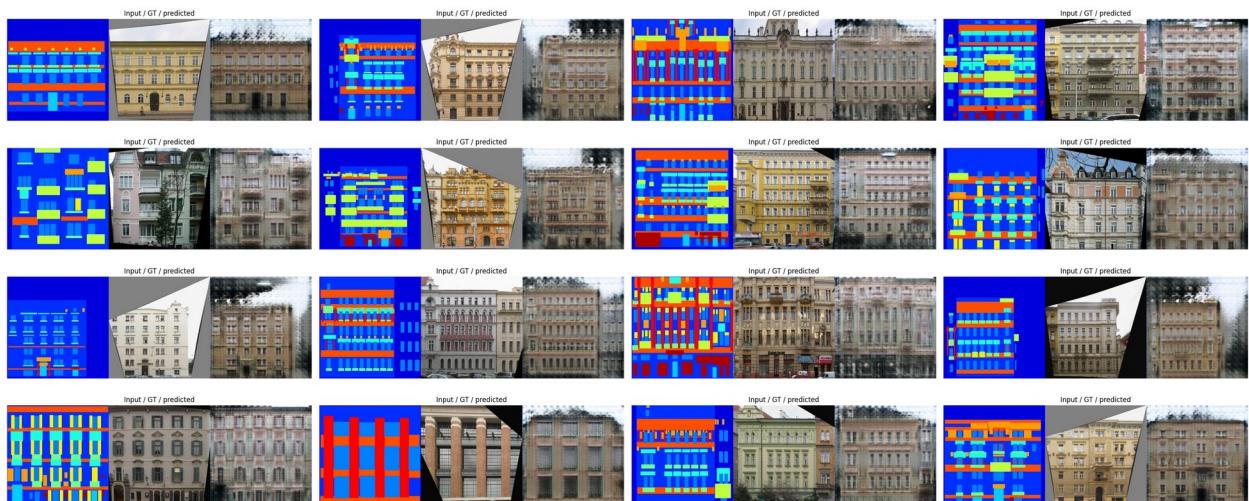
37% |  | 7500/20001 [18:22<29:52, 6.97it/s]

```
tensor(0.0328, device='cuda:0', grad_fn=<AddBackward0>)
tensor(36.7031, device='cuda:0', grad_fn=<AddBackward0>)
```



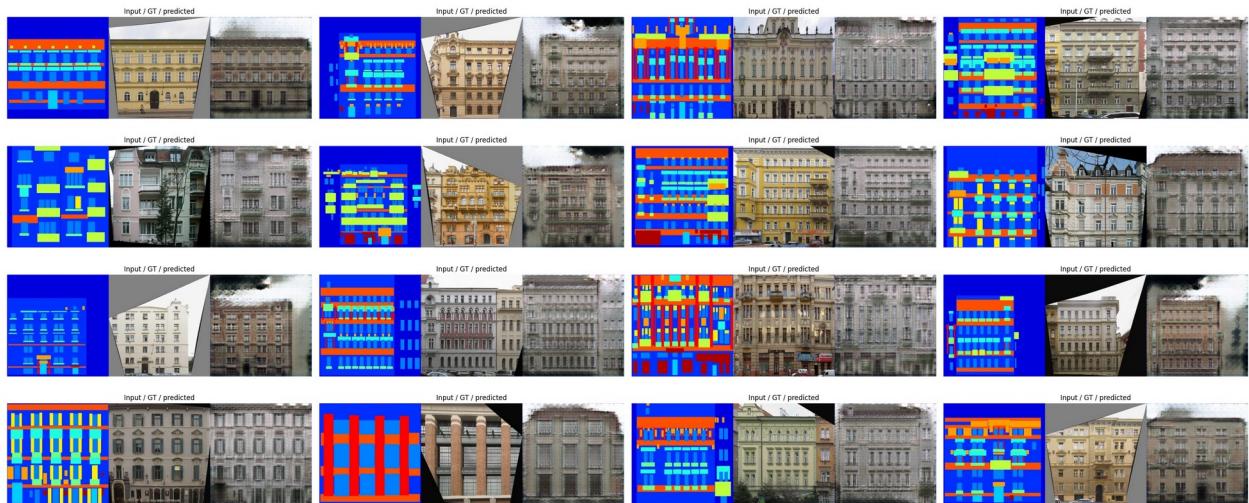
50% | [██████] | 10000/20001 [24:30<23:51, 6.99it/s]

```
tensor(0.0038, device='cuda:0', grad_fn=<AddBackward0>)
tensor(38.2071, device='cuda:0', grad_fn=<AddBackward0>)
```



62% | [██████] | 12500/20001 [30:38<18:17, 6.84it/s]

```
tensor(0.0115, device='cuda:0', grad_fn=<AddBackward0>)
tensor(34.7392, device='cuda:0', grad_fn=<AddBackward0>)
```



63% | [██████████] | 12649/20001 [31:03<17:43, 6.92it/s]

(Optional)

Combine the WGAN-GP loss with the pix2pix objective.

pass